李宏毅 (Hung-yi Lee) · HYLEE | Machine Learning (2021)

HYLEE(2021)· 课程资料包 @ShowMeAI

视频
中英双语字幕

课件
一键打包下载

笔记
官方笔记翻译

代码
作业项目解析

视频 · B 站 [扫码或点击链接]
https://www.bilibili.com/video/BV1fM4y137M4

课件 & 代码 · 博客 [扫码或点击链接]
http://blog.showmeai.tech/ntu-hylee-ml

机器学习 Auto-encoder 生成式对抗网络 学习率
深度学习 卷积神经网络 GAN 自监督 自注意力机制
批次标准化 神经网络压缩 强化学习 元学习 Transformer

Awesome AI Courses Notes Cheatsheets 是 **ShowMeAI** 资料库的分支系列，覆盖最具知名度的 **TOP50+** 门 AI 课程，旨在为读者和学习者提供一整套高品质中文学习笔记和速查表。

点击课程名称，跳转至课程**资料包**页面，**一键下载**课程全部资料！

机器学习	深度学习	自然语言处理	计算机视觉
Stanford · CS229	Stanford · CS230	Stanford · CS224n	Stanford · CS231n
# Awesome AI Courses Notes Cheatsheets· 持续更新中			
知识图谱	**图机器学习**	**深度强化学习**	**自动驾驶**
Stanford · CS520	Stanford · CS224W	UCBerkeley · CS285	MIT · 6.S094

微信公众号

资料下载方式 2：扫码点击底部菜单栏
称为 **AI 内容创作者？** 回复 [添砖加瓦]

Meta Learning: Learn to learn

Hung-yi Lee

What does “meta” mean? meta-X = X about X

Source of image: https://medium.com/intuitionmachine/the-brute-force-method-of-deep-learning-innovation-58b497323ae5 (Denny Britz's graphic)

這門課的作業在做甚麼？

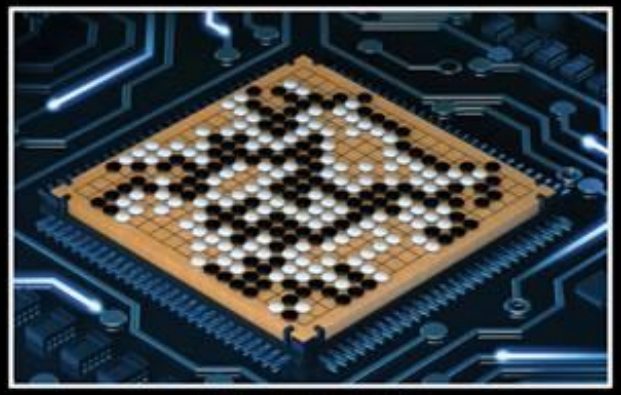

朋友覺得我在

我媽覺得我在

大眾覺得我在

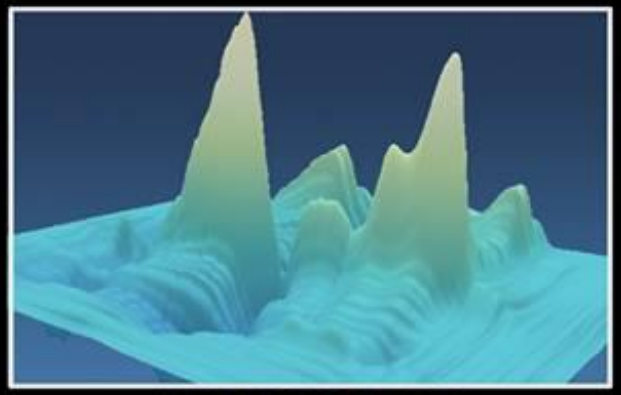

指導教授覺得我在

我以為我在

事實上我在

感謝 沈昇勳 同學提供圖檔

Industry

Using 1000 GPUs to try 1000 sets of hyperparameters

Academia

“Telepathize” (通靈) a set of good hyperparameters

Can machine automatically determine the hyperparameters?

Machine Learning 101

Machine Learning

= Looking for a function

Dog-Cat Classification

$f($ $) =$ "cat"

Step 1: Function with unknown

Step 2: Define loss function

Step 3: Optimization

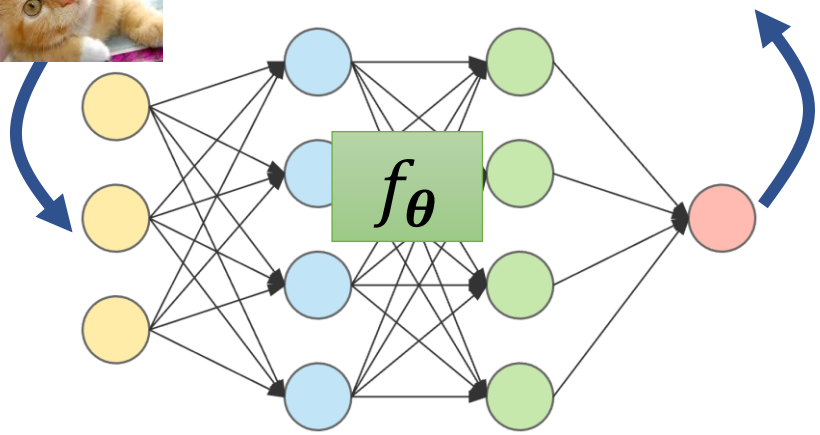

Weights and biases of neurons are unknown parameters (*learnable*).

Using $\boldsymbol{\theta}$ to represent the unknown parameters.

Machine Learning

Step 1: Function with unknown

Step 2: Define loss function

Step 3: Optimization

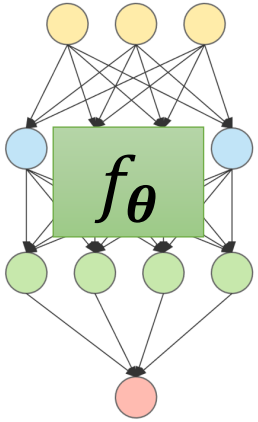

$L(\theta)$

$$L(\theta) = \sum_{k=1}^{K} e_k$$

Training Examples

f_θ f_θ

cat dog cat dog

Cross-entropy e_1 e_2

cat dog cat dog

Ground Truth

Machine Learning 101

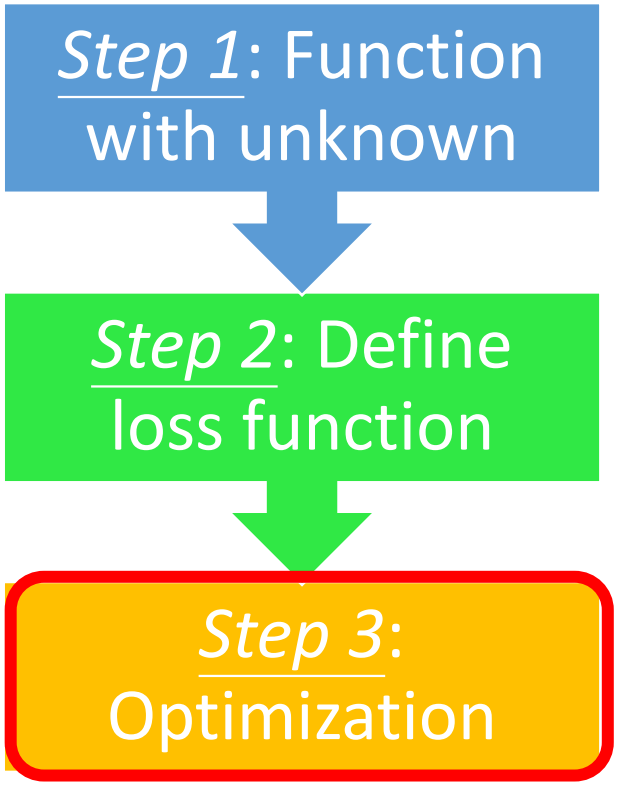

loss: $L(\boldsymbol{\theta}) = \sum_{k=1}^{K} e_k$ sum over training examples

$$\boldsymbol{\theta}^* = arg \min_{\boldsymbol{\theta}} L(\boldsymbol{\theta})$$

done by gradient descent

$f_{\boldsymbol{\theta}^*}$ is the function learned by learning algorithm from data

Introduction of
Meta Learning

What is Meta Learning?

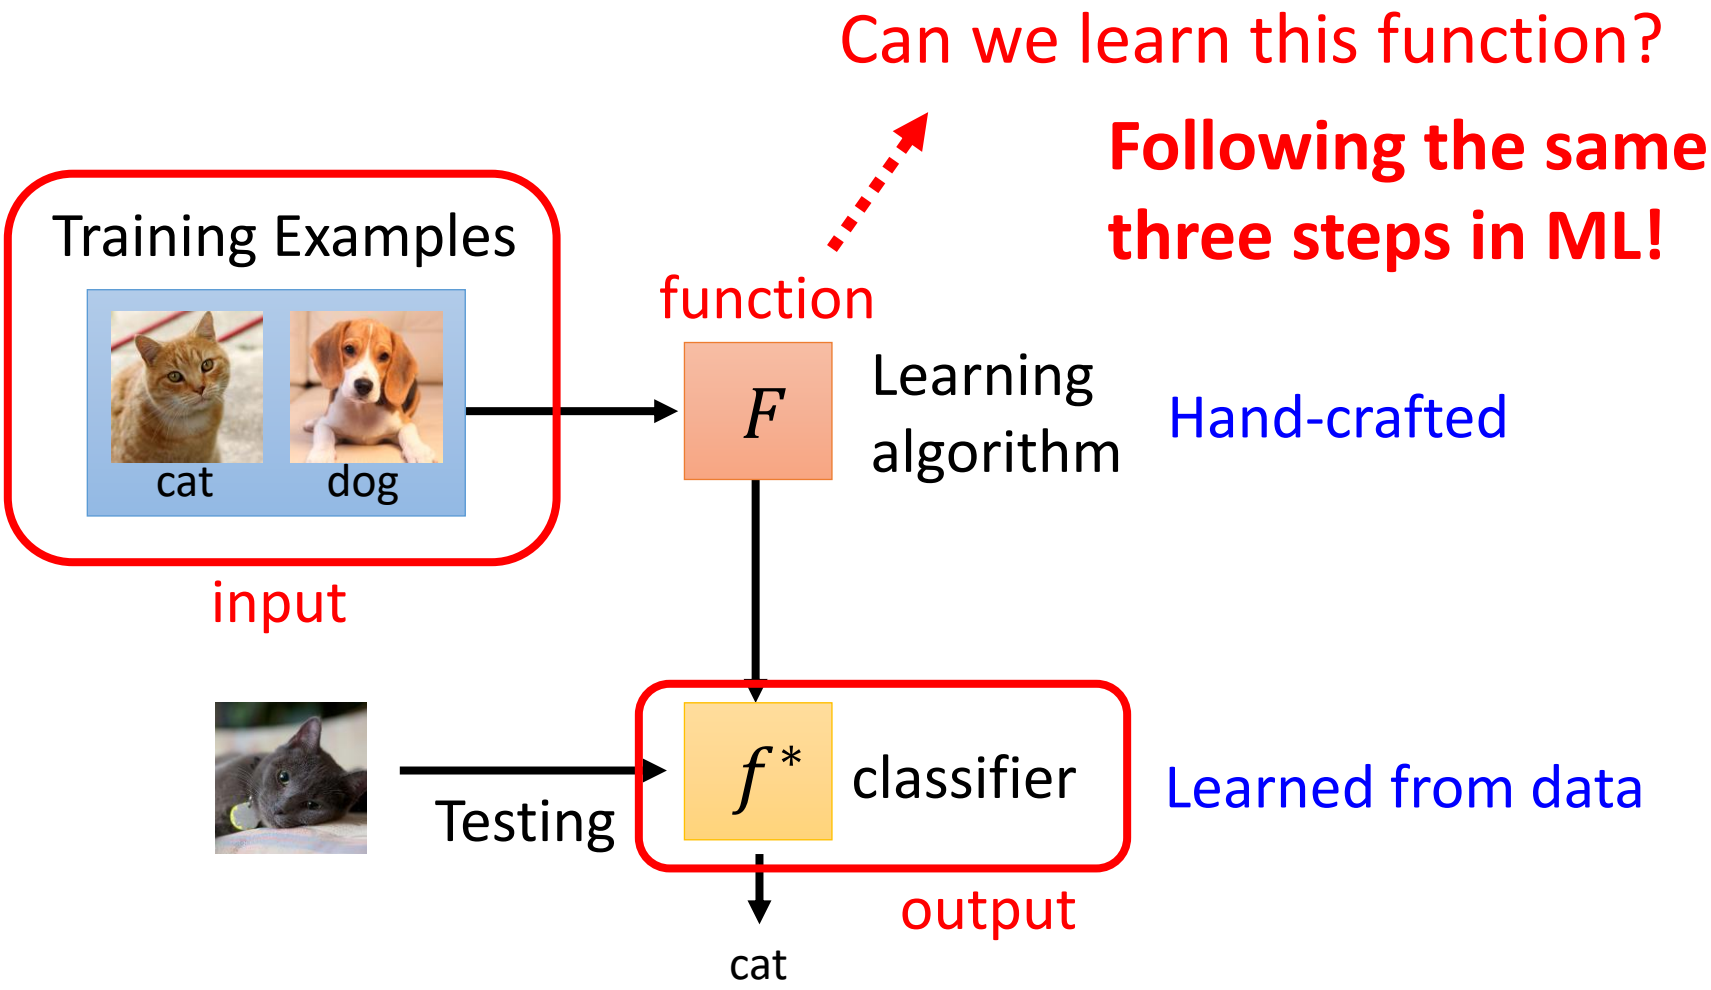

Meta Learning – Step 1

- What is ***learnable*** in a learning algorithm?

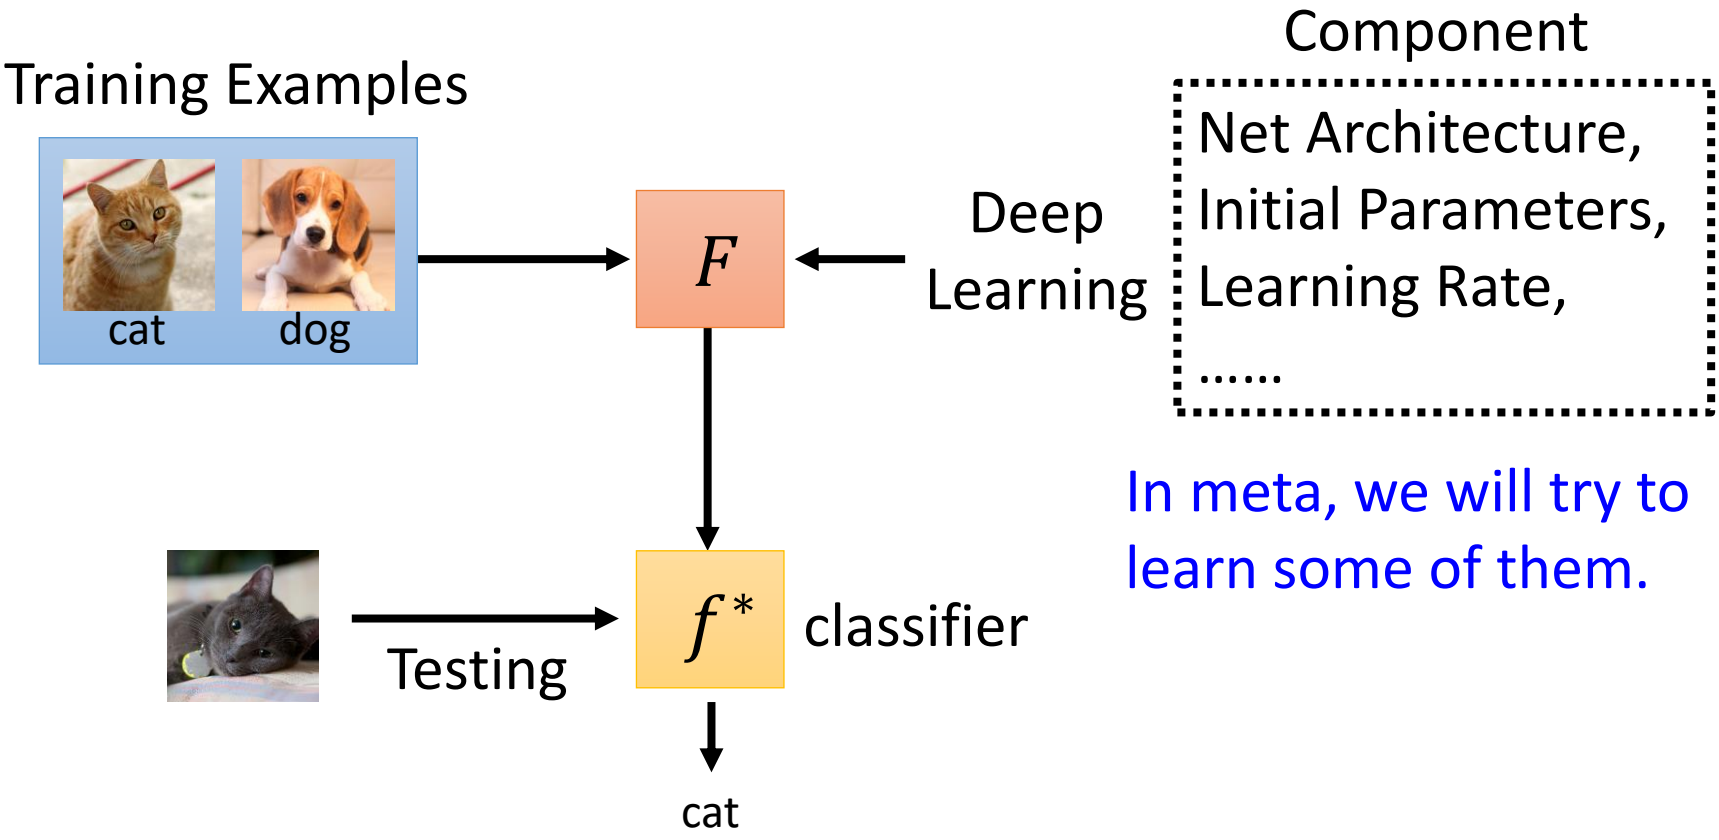

Meta Learning – Step 1

- What is ***learnable*** in a learning algorithm?

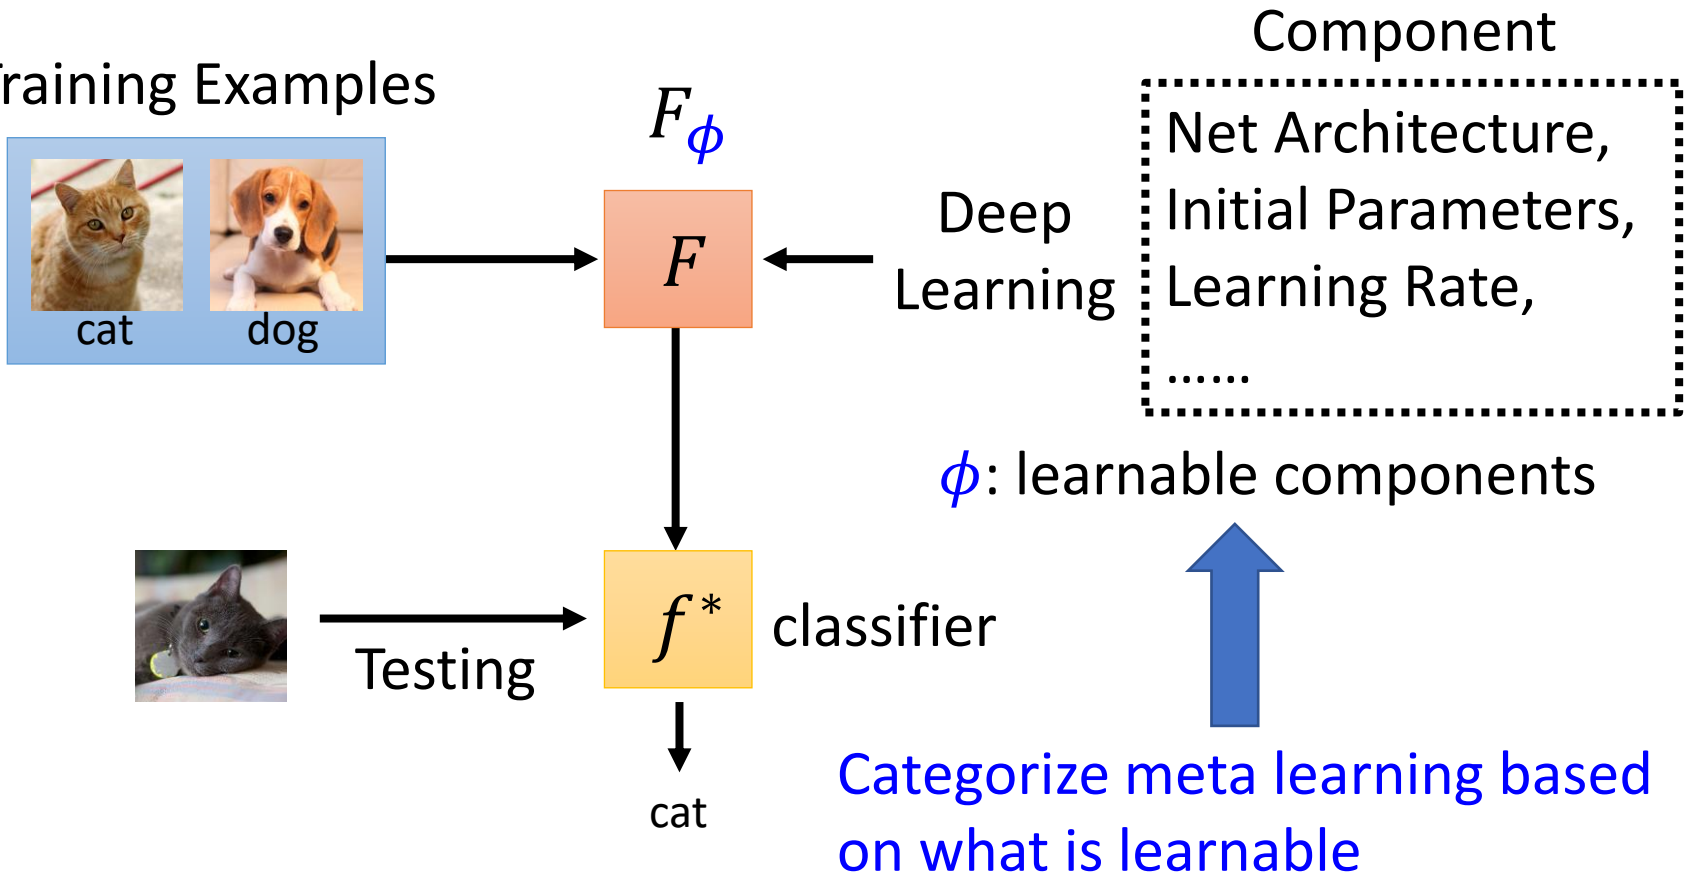

Meta Learning – Step 2

- Define ***loss function*** for ***learning algorithm*** F_ϕ

$L(\phi)$

$L(\phi)$ ↓

$L(\phi)$ ↑

Training Tasks

Task 1

Apple & Orange

Train

apple

orange

Test

apple

orange

Task 2

Car & Bike

Train

bike

car

Test

bike

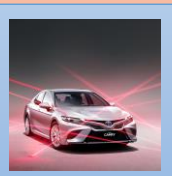
car

Meta Learning – Step 2

Task 1

Training Examples

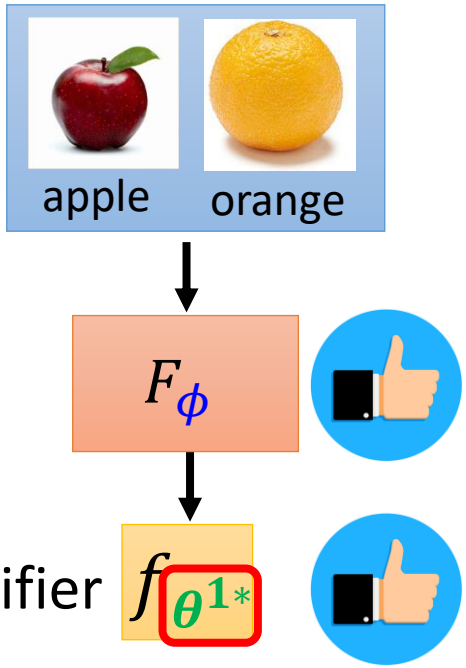

How to define $L(\phi)$

$L(\phi)$ ↓

θ^{1*}: parameters of the classifier learned by F_ϕ using the training examples of task 1

Meta Learning – Step 2

Task 1

Training Examples

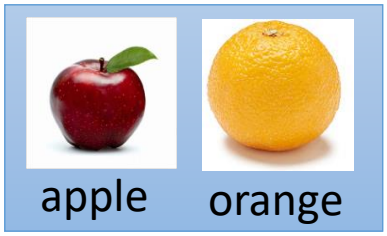

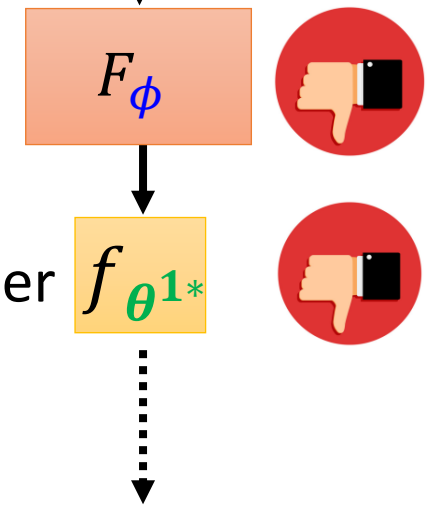

How to define $L(\phi)$

$L(\phi)$ ↑

How can we know a classifier is good or bad?

Evaluate the classifier on testing set

Meta Learning – Step 2

Task 1

Training Examples

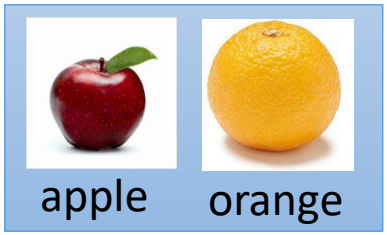

F_ϕ

Testing Examples

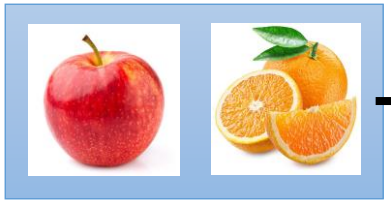

$f_{\theta^{1*}}$

prediction

l^1

Compute *difference*

Testing Examples

$f_{\theta^{1*}}$ $f_{\theta^{1*}}$

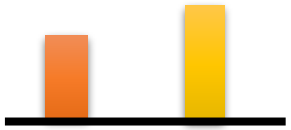

Cross-entropy Cross-entropy

Ground Truth

Meta Learning – Step 2

Task 1

Training Examples

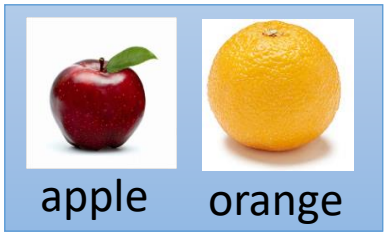

F_ϕ

Testing Examples

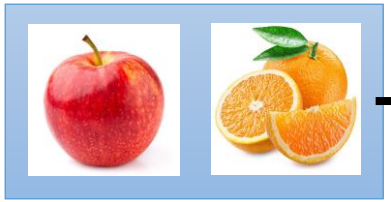

$f_{\theta^{1*}}$

prediction

l^1

Compute *difference*

Testing Examples

$f_{\theta^{1*}}$ $f_{\theta^{1*}}$

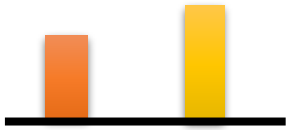

Cross-entropy Cross-entropy

Ground Truth

Meta Learning – Step 2

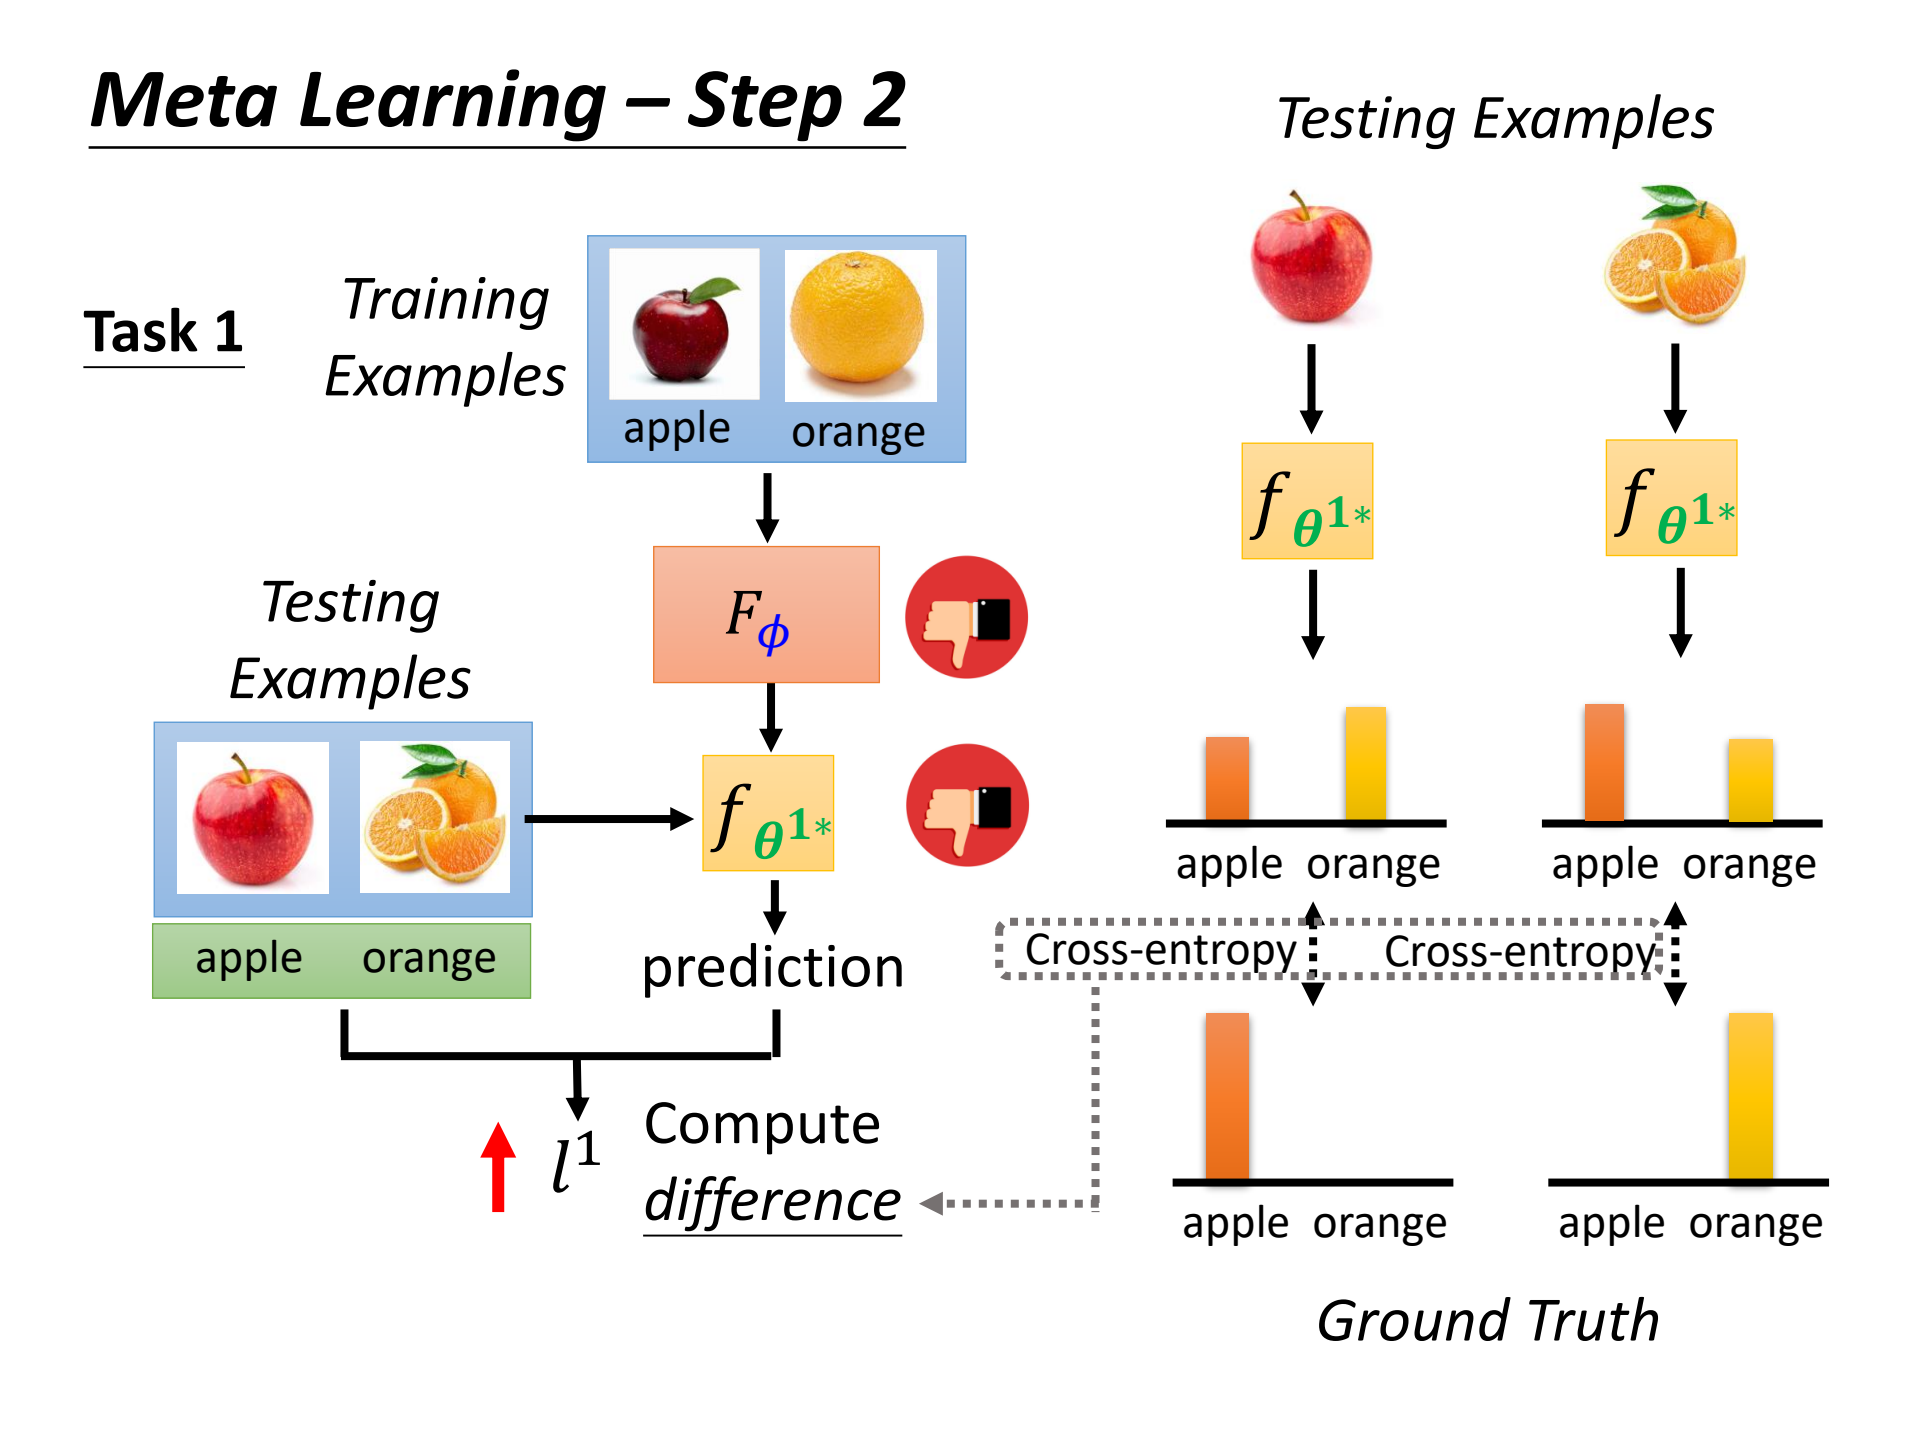

Meta Learning – Step 2

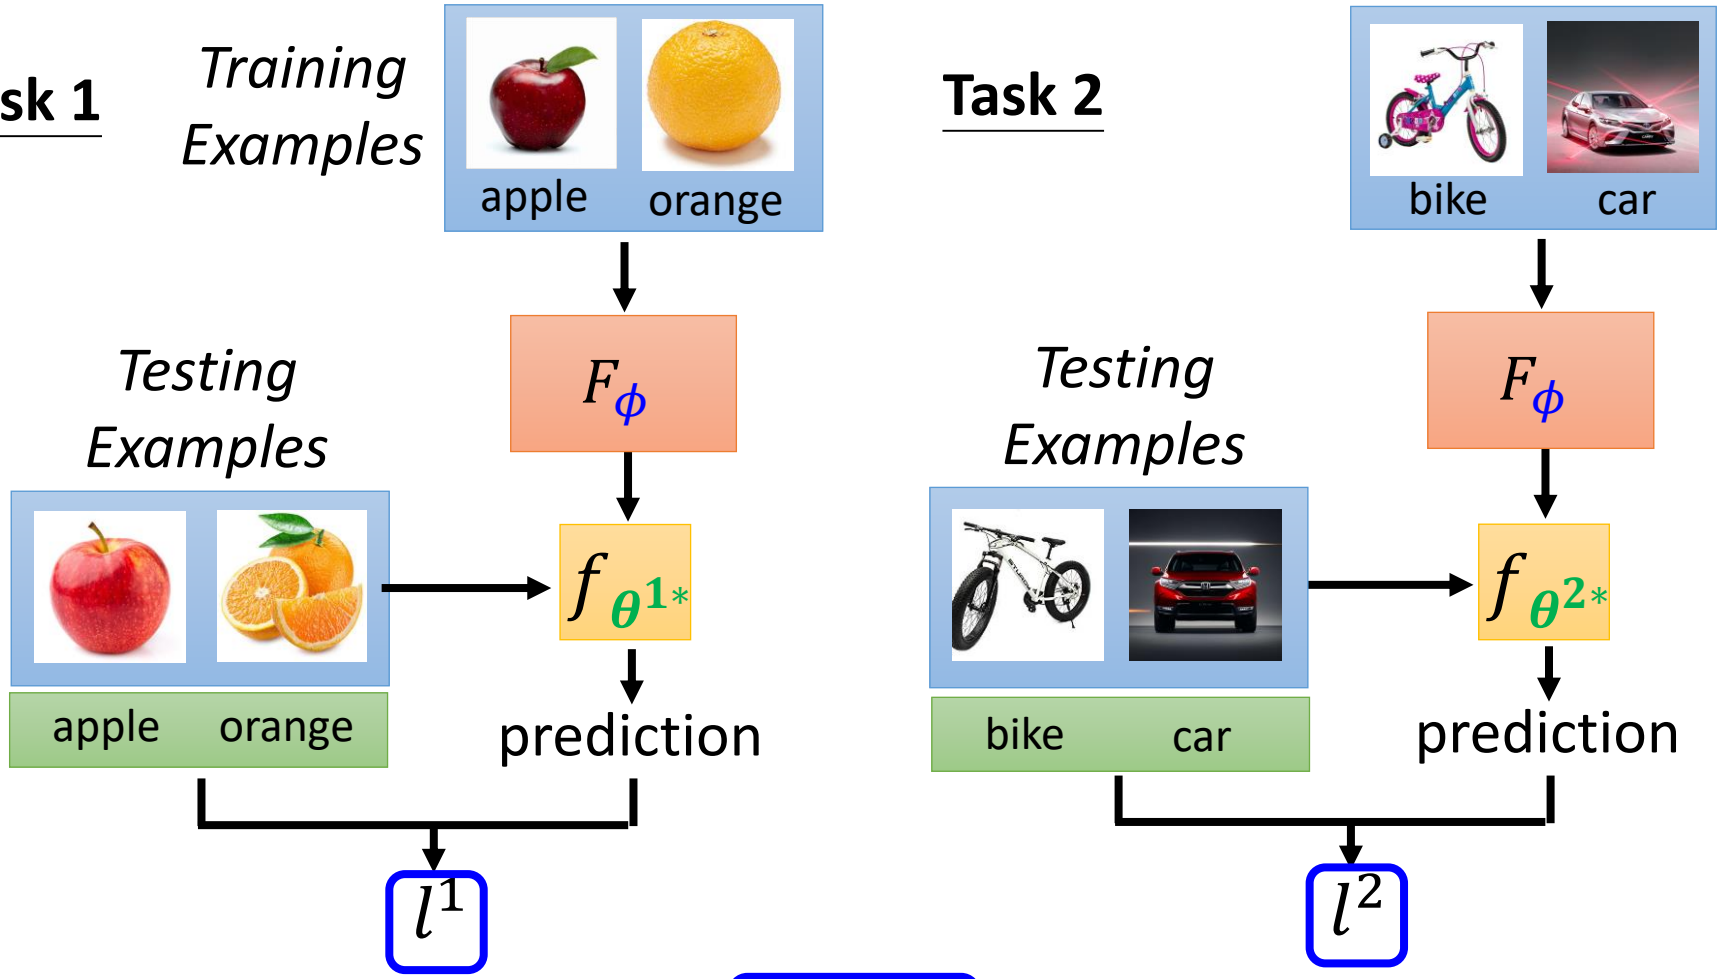

Total loss: $L(\phi) = l^1 + l^2$ (sum over all the training tasks)

Meta Learning – Step 2

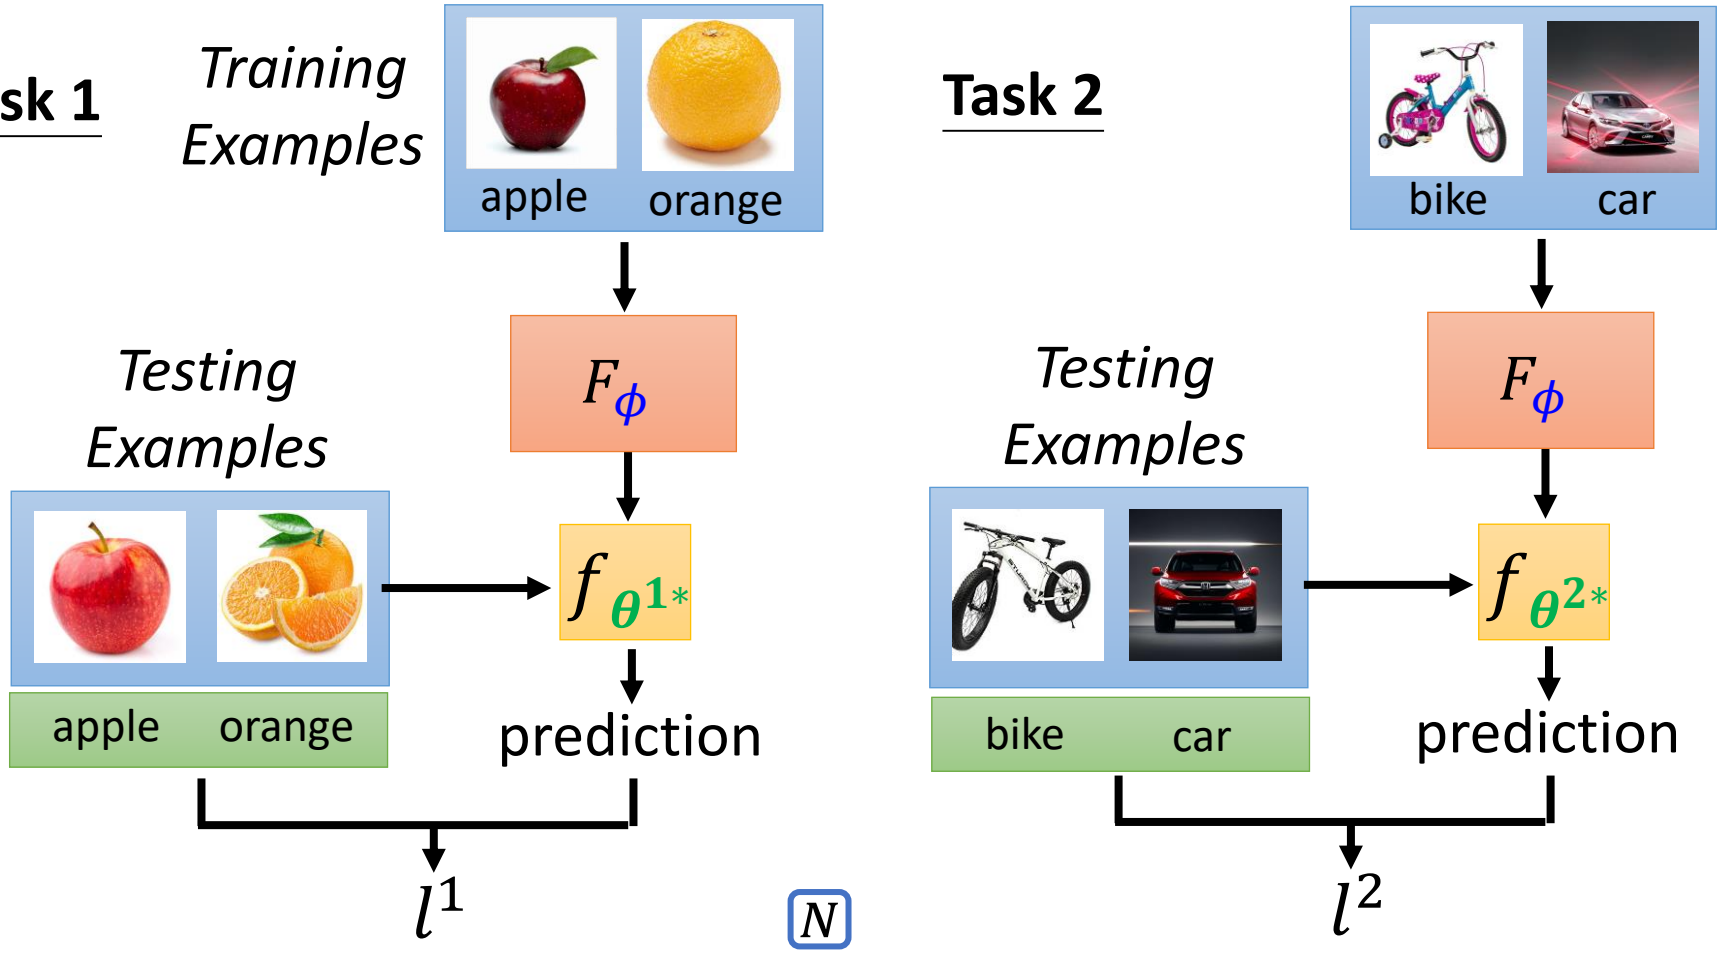

Total loss: $L(\phi) = \sum_{n=1}^{N} l^n$ (N is the number of the training tasks)

Meta Learning – Step 2

Task 1

In typical ML, you compute the loss based on training examples

In meta, you compute the loss based on testing examples

Testing Examples

$f_{\theta^{1*}}$ $f_{\theta^{1*}}$

Hold on! You use testing examples during training???

apple orange

prediction

l^1

Compute *difference*

apple orange apple orange

Ground Truth

Meta Learning – Step 2

Task 1

In typical ML, you compute the loss based on training examples

In meta, you compute the loss based on testing examples of training tasks.

Testing Examples

$f_{\theta^{1*}}$ $f_{\theta^{1*}}$

apple orange → prediction

l^1 → Compute *difference*

Meta Learning – Step 3

- Loss function for learning algorithm $L(\phi) = \sum_{n=1}^{N} l^n$
- Find ϕ that can minimize $L(\phi)$ $\phi^* = arg \min_{\phi} L(\phi)$
- Using the optimization approach you know

  If you know how to compute $\partial L(\phi)/\partial \phi$

  Gradient descent is your friend.

  What if $L(\phi)$ is not differentiable?

  Reinforcement Learning / Evolutionary Algorithm

Now we have a learned "learning algorithm" F_{ϕ^*}

Framework

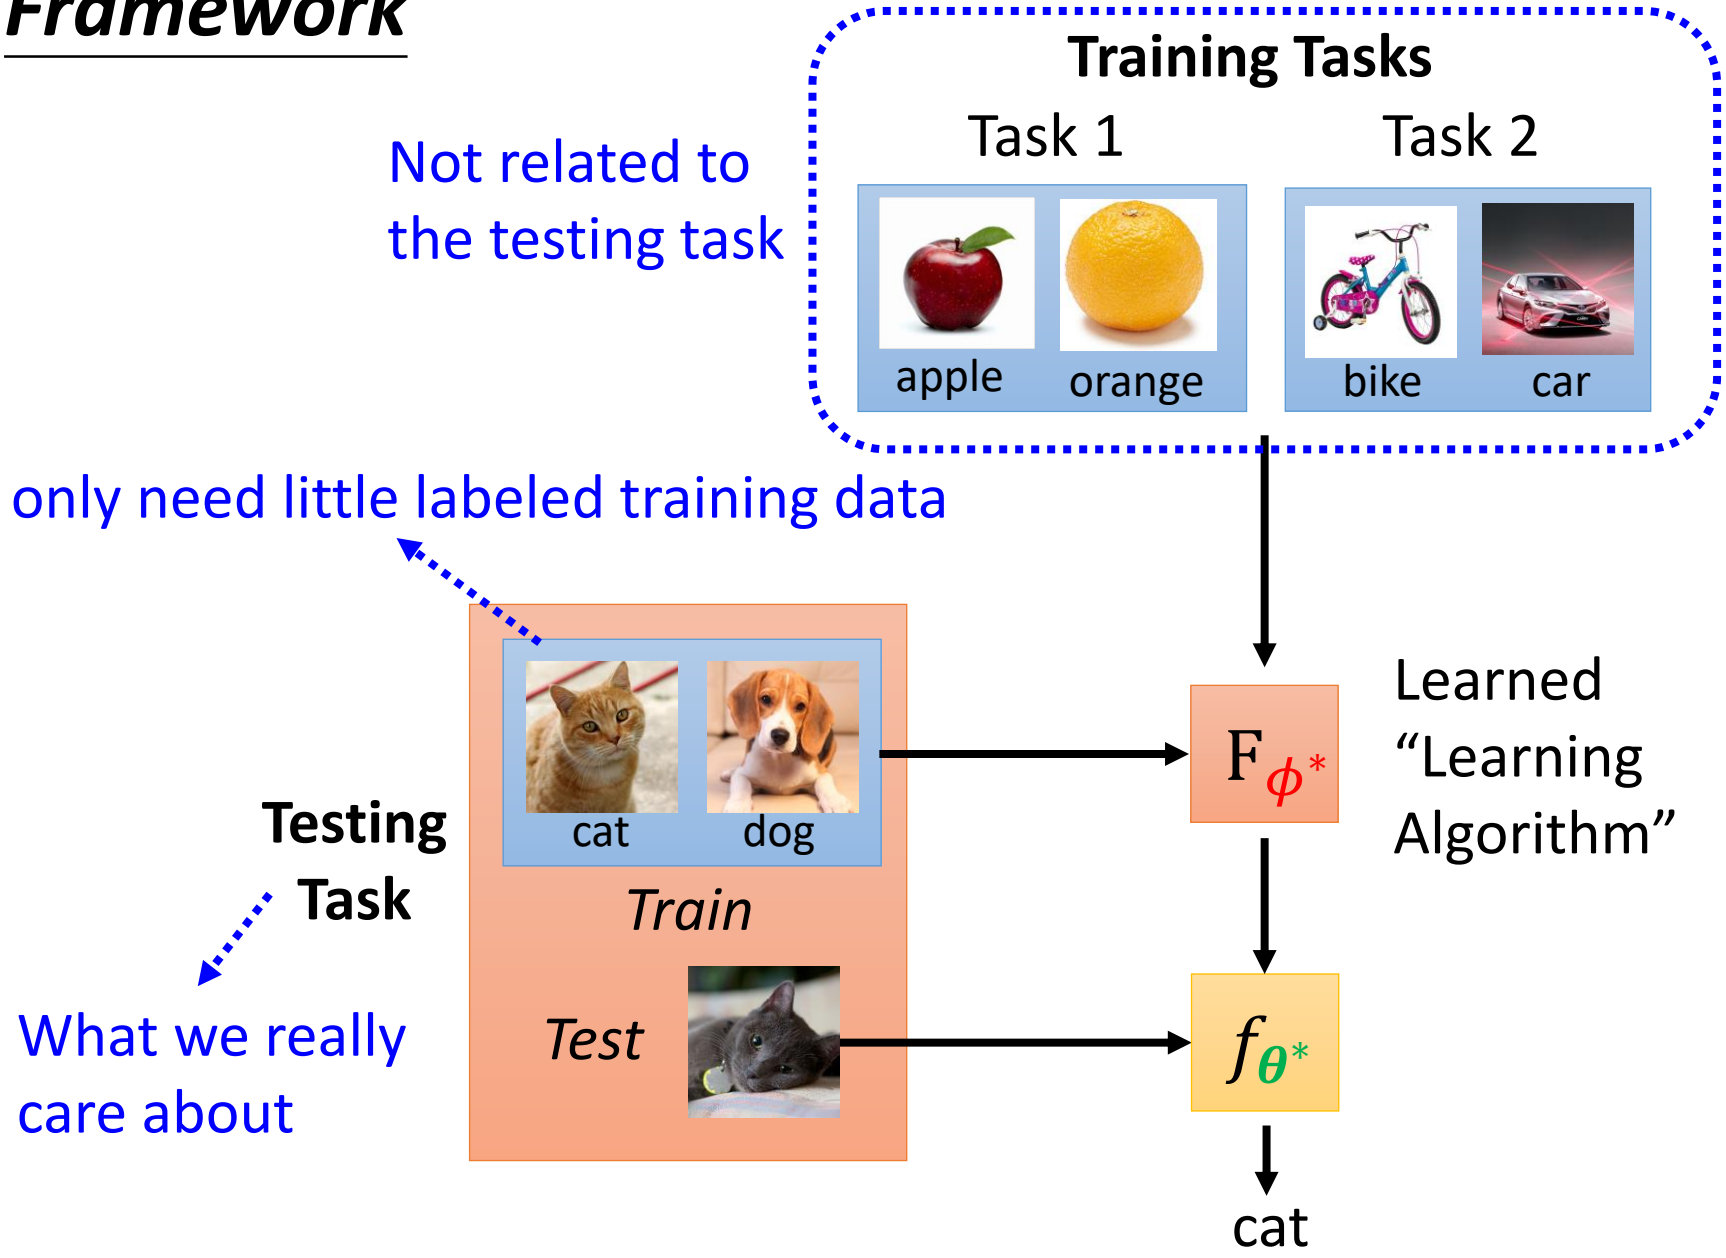
Training Tasks
Task 1
Task 2
Not related to the testing task
apple
orange
bike
car
only need little labeled training data
cat
dog
Train
Testing Task
Test
What we really care about
F_{ϕ^*}
f_{θ^*}
Learned "Learning Algorithm"
cat

ML v.s. Meta

Goal

Machine Learning ≈ find a function f

Dog-Cat Classification

$f($ $) =$ "cat"

Meta Learning

≈ find a function F that finds a function f

Learning Algorithm

$F($ 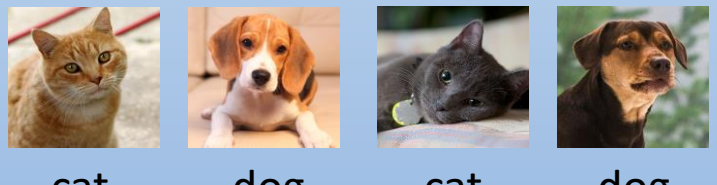

$) = f$

Training Data

Machine Learning

One task

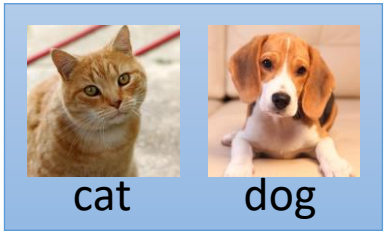

Train

Meta Learning

Training tasks

Task 1

Apple & Orange

Train

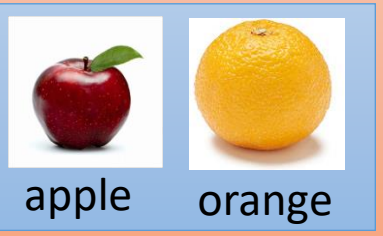

Test

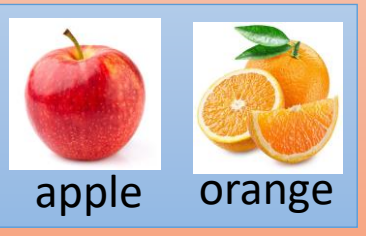

Task 2

Car & Bike

Train

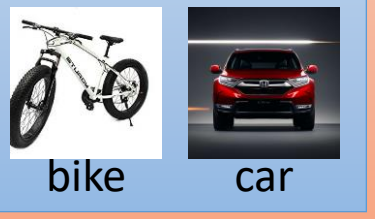

Test

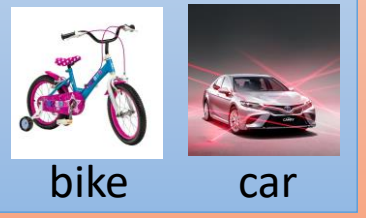

Support set

Query set

(in the literature of "*learning to compare*")

Machine Learning

Within-task Training

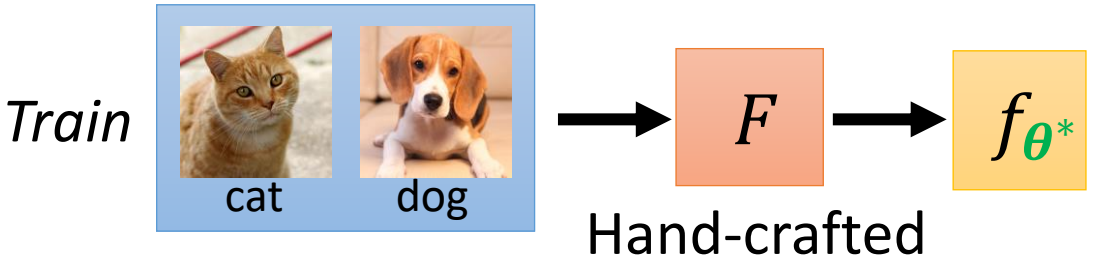

Meta Learning

Training Tasks

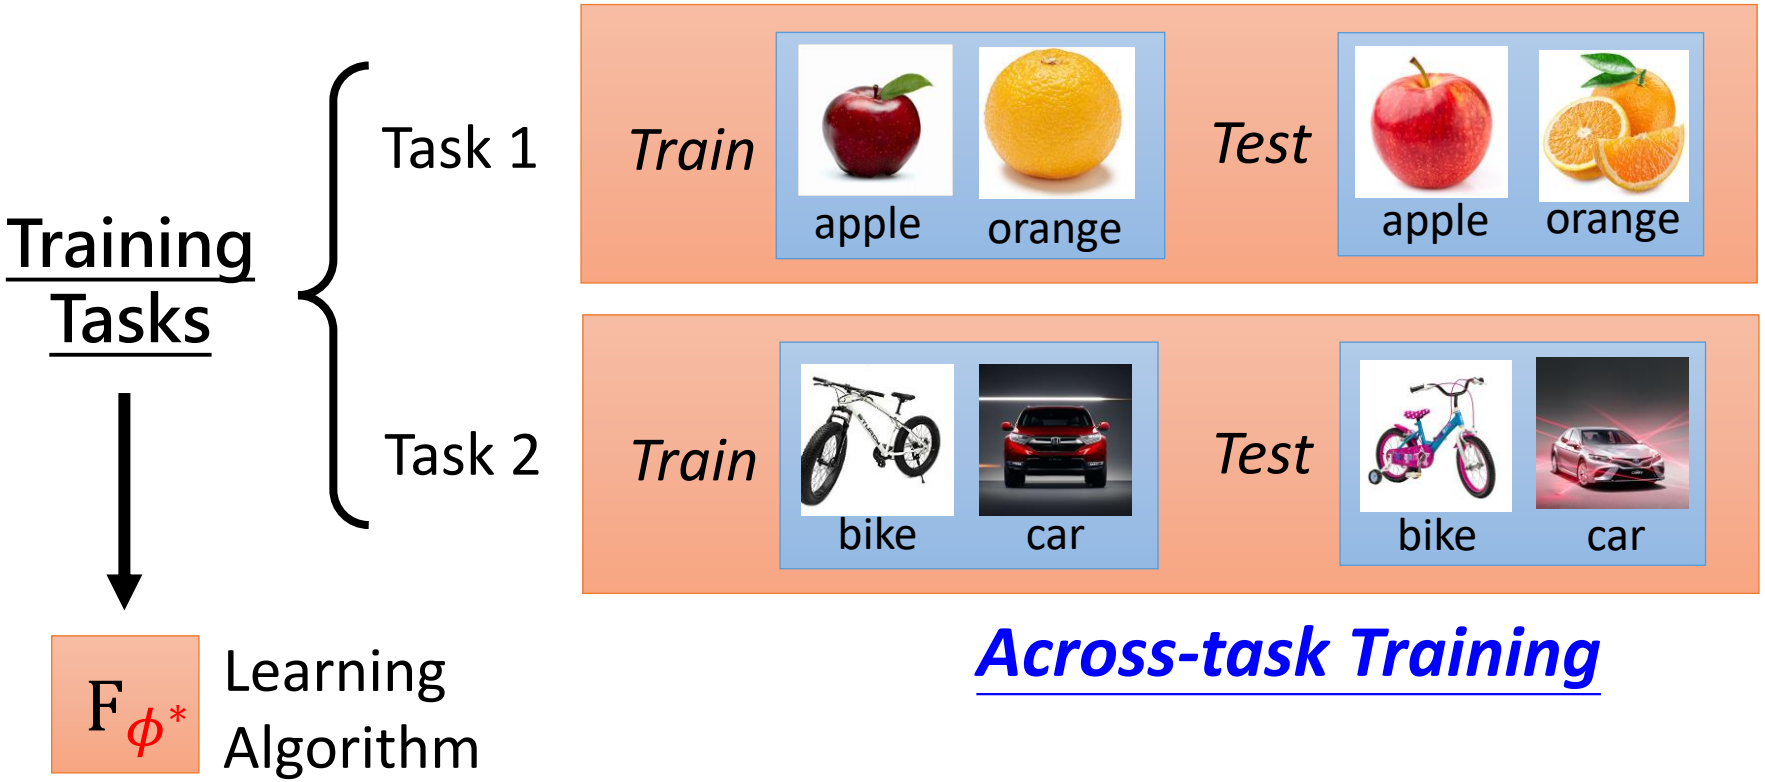

Across-task Training

Machine Learning

Training Examples

Test

Within-task Testing

f_{θ^*}

cat

Meta Learning

Training Tasks

Testing Task

Train

Test

Within-task Training

F_{ϕ^*}

Learned "Learning Algorithm"

Within-task Testing

f_{θ^*}

cat

Episode

Across-task Testing

Loss

Machine Learning

$$L(\boldsymbol{\theta}) = \sum_{k=1}^{K} e_k$$

Sum over training examples in one task

Meta Learning

$$L(\phi) = \sum_{n=1}^{N} l^n$$

Sum over testing examples in one task

Sum over training tasks

$$L(\phi) = \sum_{n=1}^{N} l^n$$

If your optimization method needs to compute $L(\phi)$

Outer Loop in "*Learning to initialize*"

Across-task training includes **within-task training and testing**

Training Examples

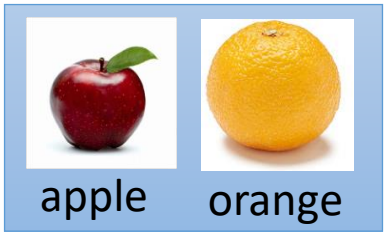

F_ϕ

Testing Examples

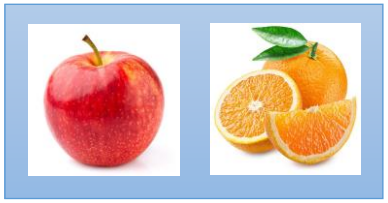

f_{θ^*}

prediction

Inner Loop in "*Learning to initialize*"

Within-task Training

Within-task Testing

l^1 ..... To compute the loss

Meta Learning v.s ML

- What you know about ML can usually apply to meta learning
  - Overfitting on training tasks
  - Get more training tasks to improve performance
  - Task augmentation
  - There are also hyperparameters when learning a learning algorithm ……
  - Development task ☺

What is learnable in a learning algorithm?

Review: Gradient Descent

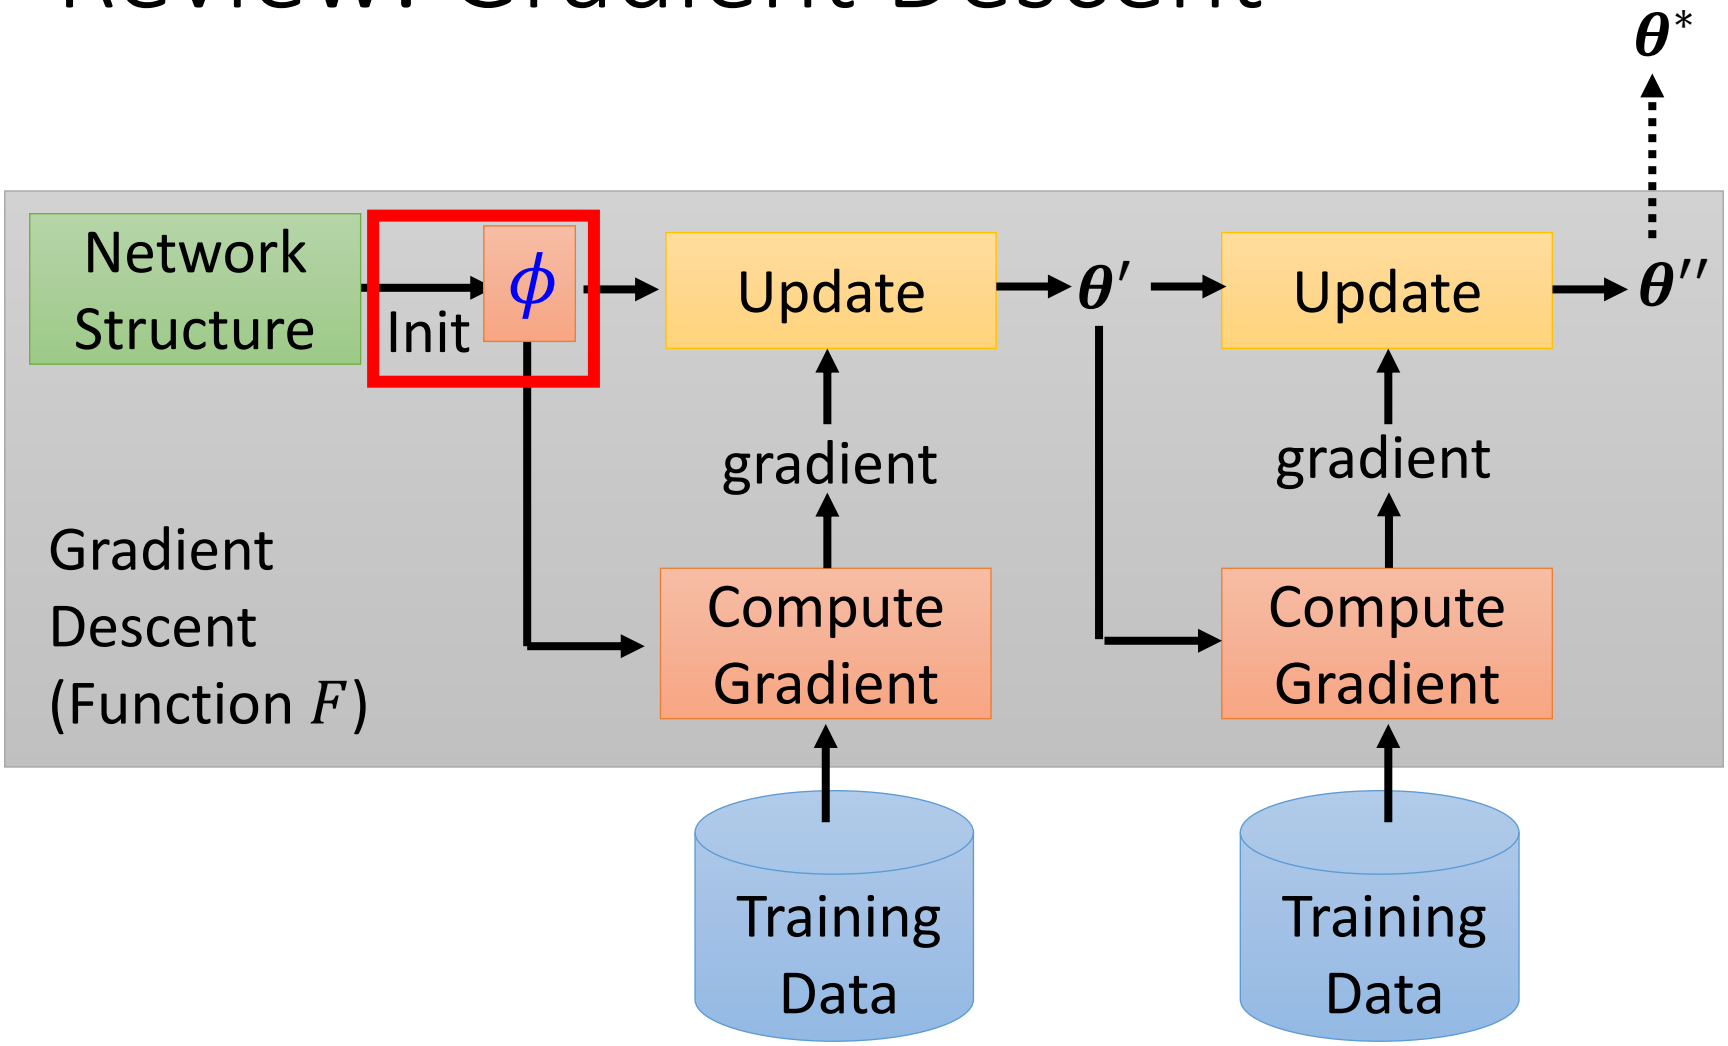

Learning to initialize

- Model-Agnostic Meta-Learning (MAML)

Mammals

Chelsea Finn, Pieter Abbeel, and Sergey Levine, “Model-Agnostic Meta-Learning for Fast Adaptation of Deep Networks”, ICML, 2017

- Reptile

https://arxiv.org/abs/1803.02999

How to train your ~~Dragon~~ MAML

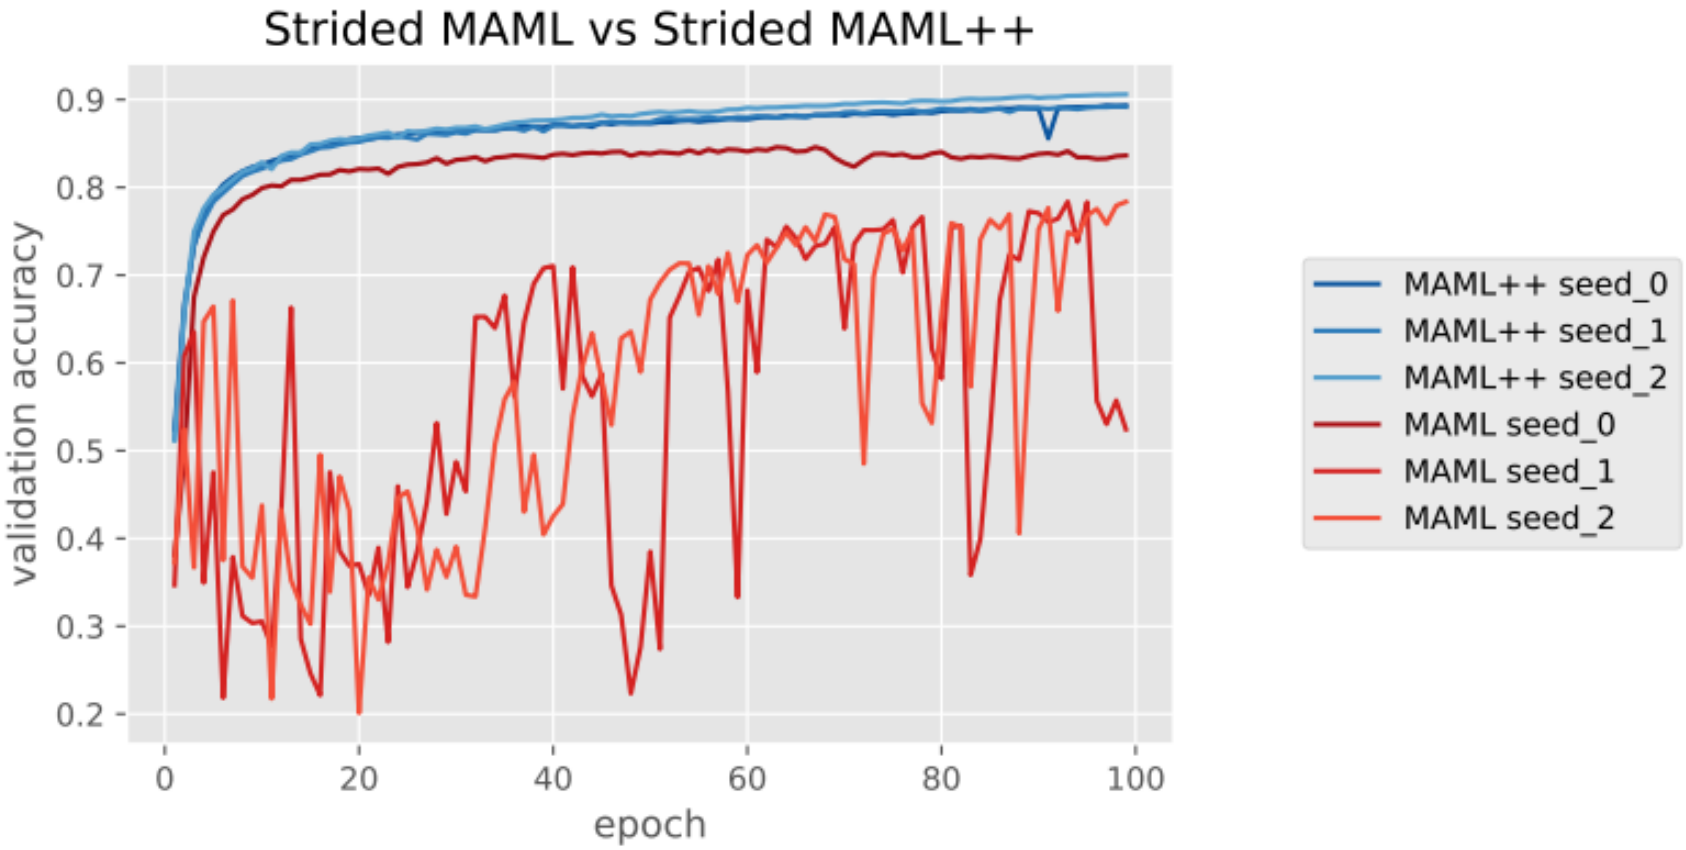

Antreas Antoniou, Harrison Edwards, Amos Storkey, How to train your MAML, ICLR, 2019

MAML

Task 1

cat dog

Task 2

cat dog

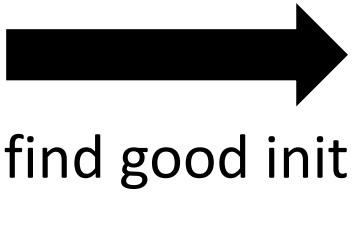

find good init

Testing Task

cat dog

Pre-training (Self-supervised Learning)

Trained by proxy tasks
(fill-in the blanks, etc.)

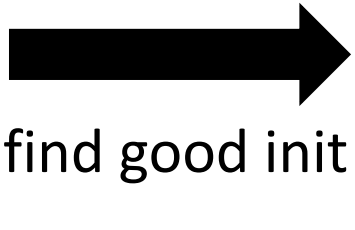

find good init

cat dog

MAML

Isn't it domain adaptation / transfer learning?

Task 1

cat dog

Task 2

cat dog

find good init

cat dog

Pre-training (more typical ways)

cat dog

cat dog

find good init

cat dog

Use data from different tasks to train a model

Also known as multi-task learning (baseline of meta)

MAML v.s. Pre-training

- https://youtu.be/vUwOA3SNb_E

影片中有防不勝防
的業配

這就是 “meta 業配”

MAML is good because ……

- ANIL (Almost No Inner Loop)

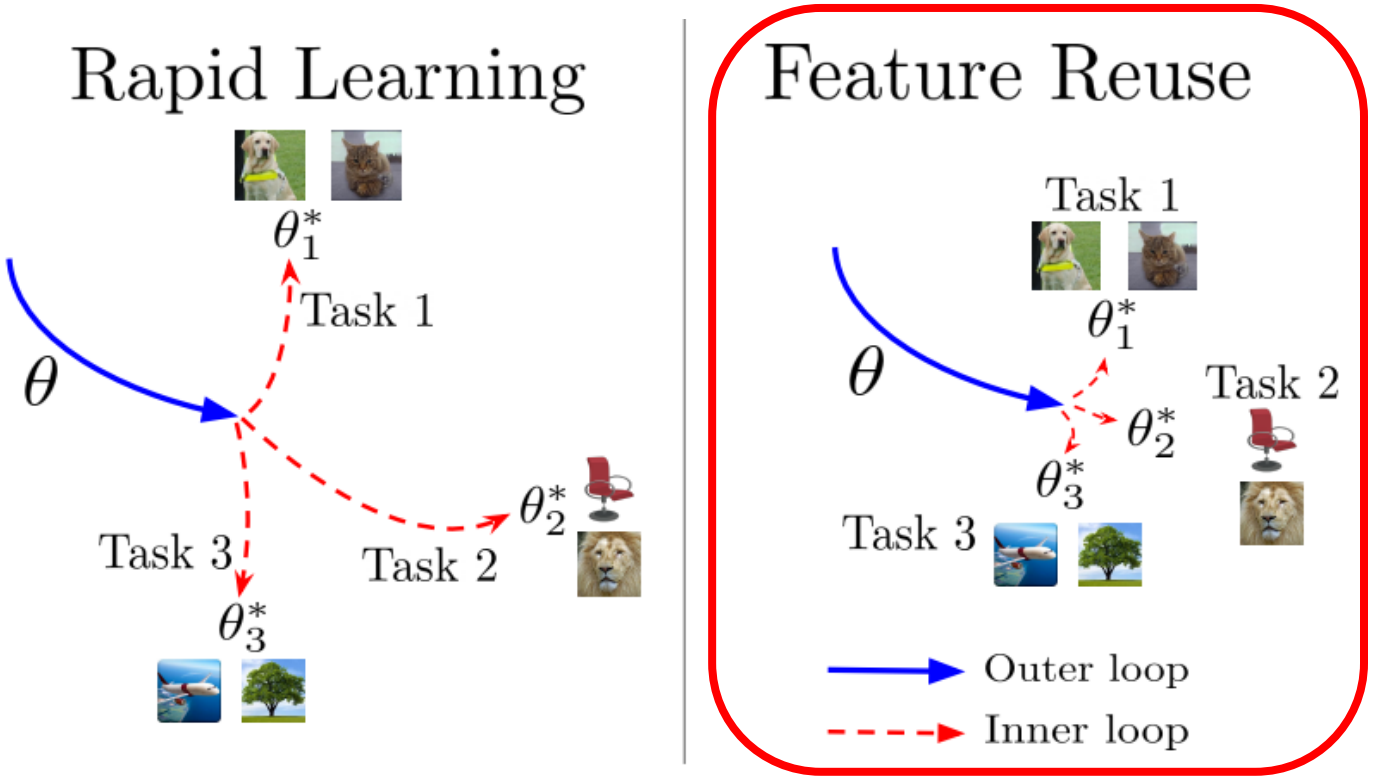

Aniruddh Raghu, Maithra Raghu, Samy Bengio, Oriol Vinyals, Rapid Learning or Feature Reuse? Towards Understanding the Effectiveness of MAML, ICLR, 2020

More about MAML

- More mathematical details behind MAML
  - https://youtu.be/mxqzGwP_Qys
- First order MAML (FOMAML)
  - https://youtu.be/3z997JhL9Oo
- Reptile
  - https://youtu.be/9jJe2AD35P8

Optimizer

Basis form: $\boldsymbol{\theta^{t+1}} \leftarrow \boldsymbol{\theta^t} - \lambda \boldsymbol{g^t}$ ϕ

Adagrad, RMSprop, NAG, Adam ……

Is the optimizer learnable?

Can be learned by MAML

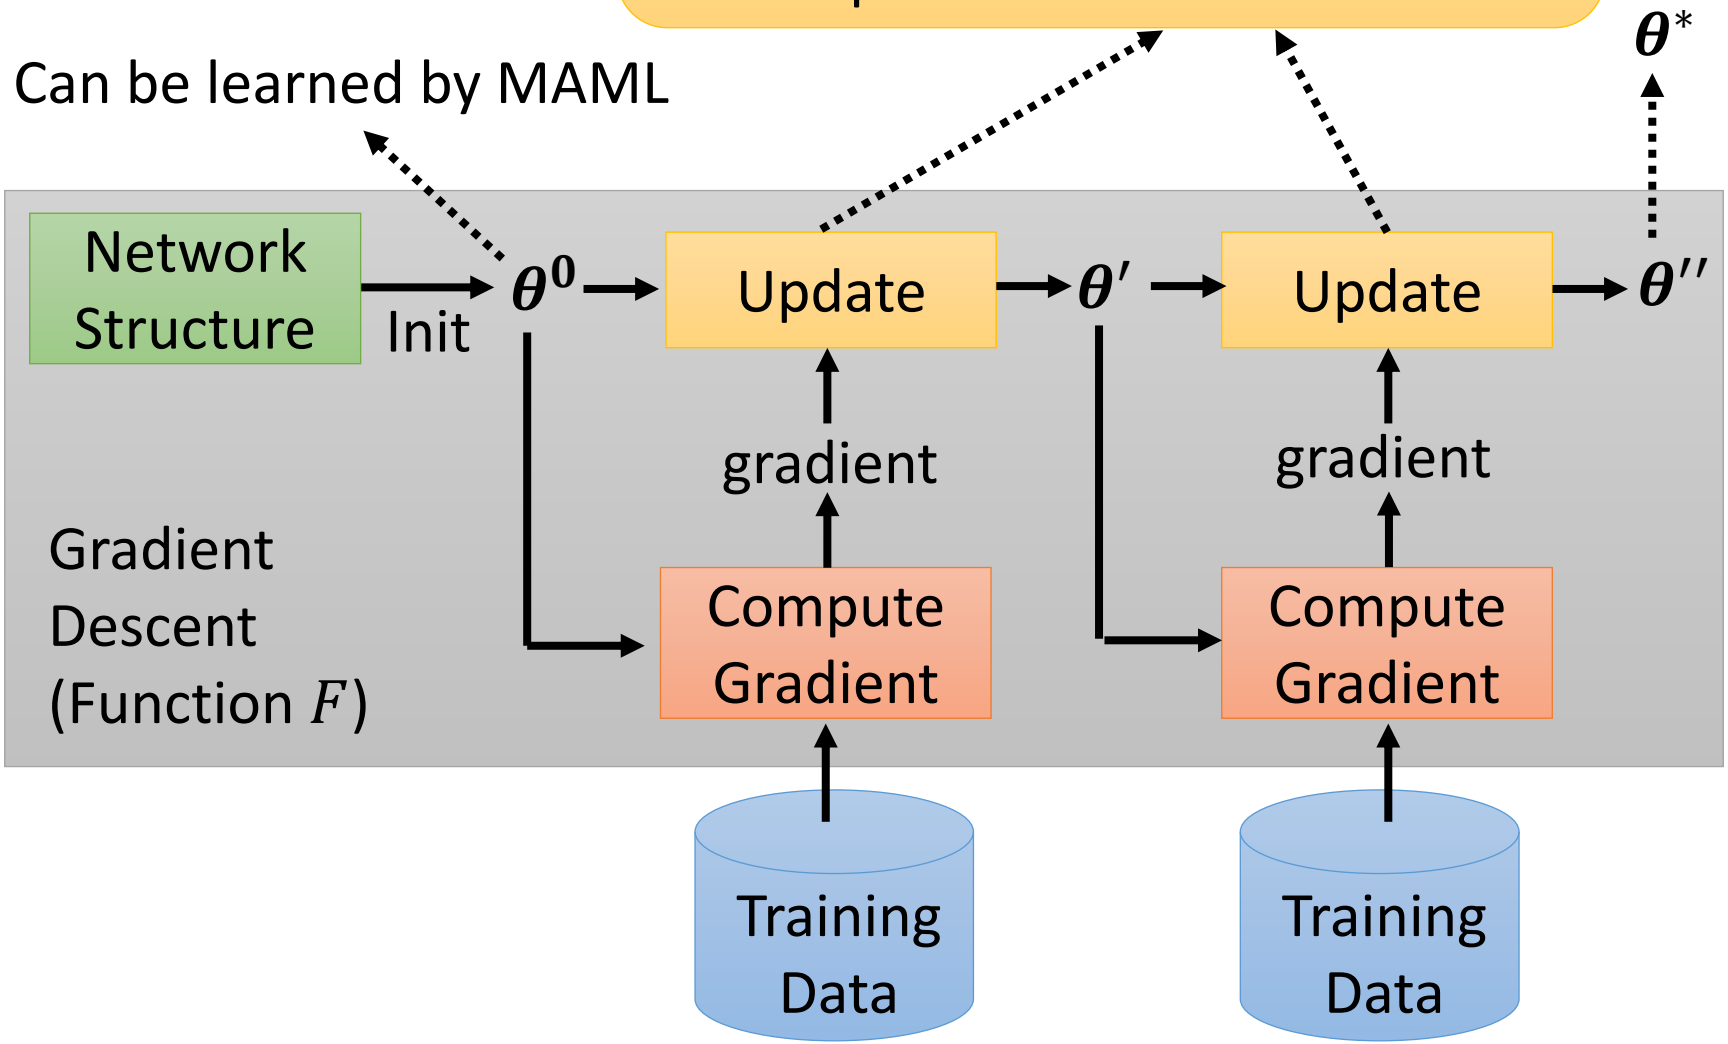

Optimizer

Marcin Andrychowicz, et al., Learning to learn by gradient descent by gradient descent, NIPS, 2016

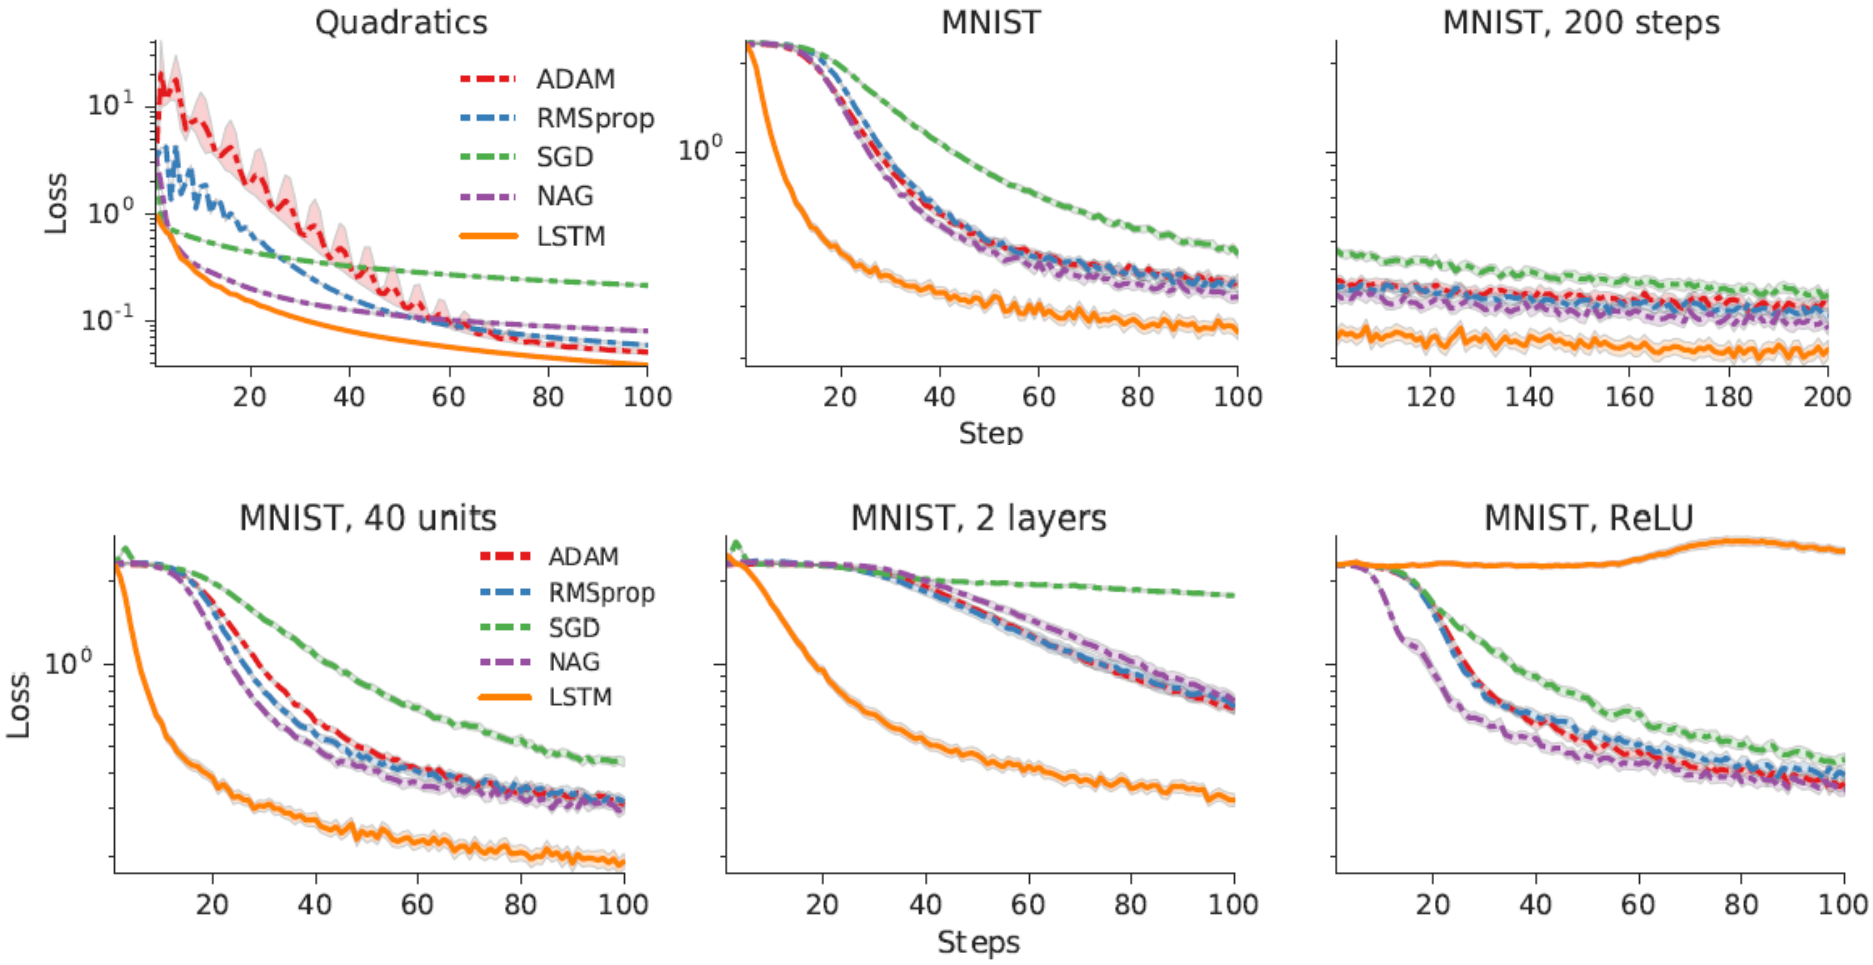

Network Architecture Search (NAS)

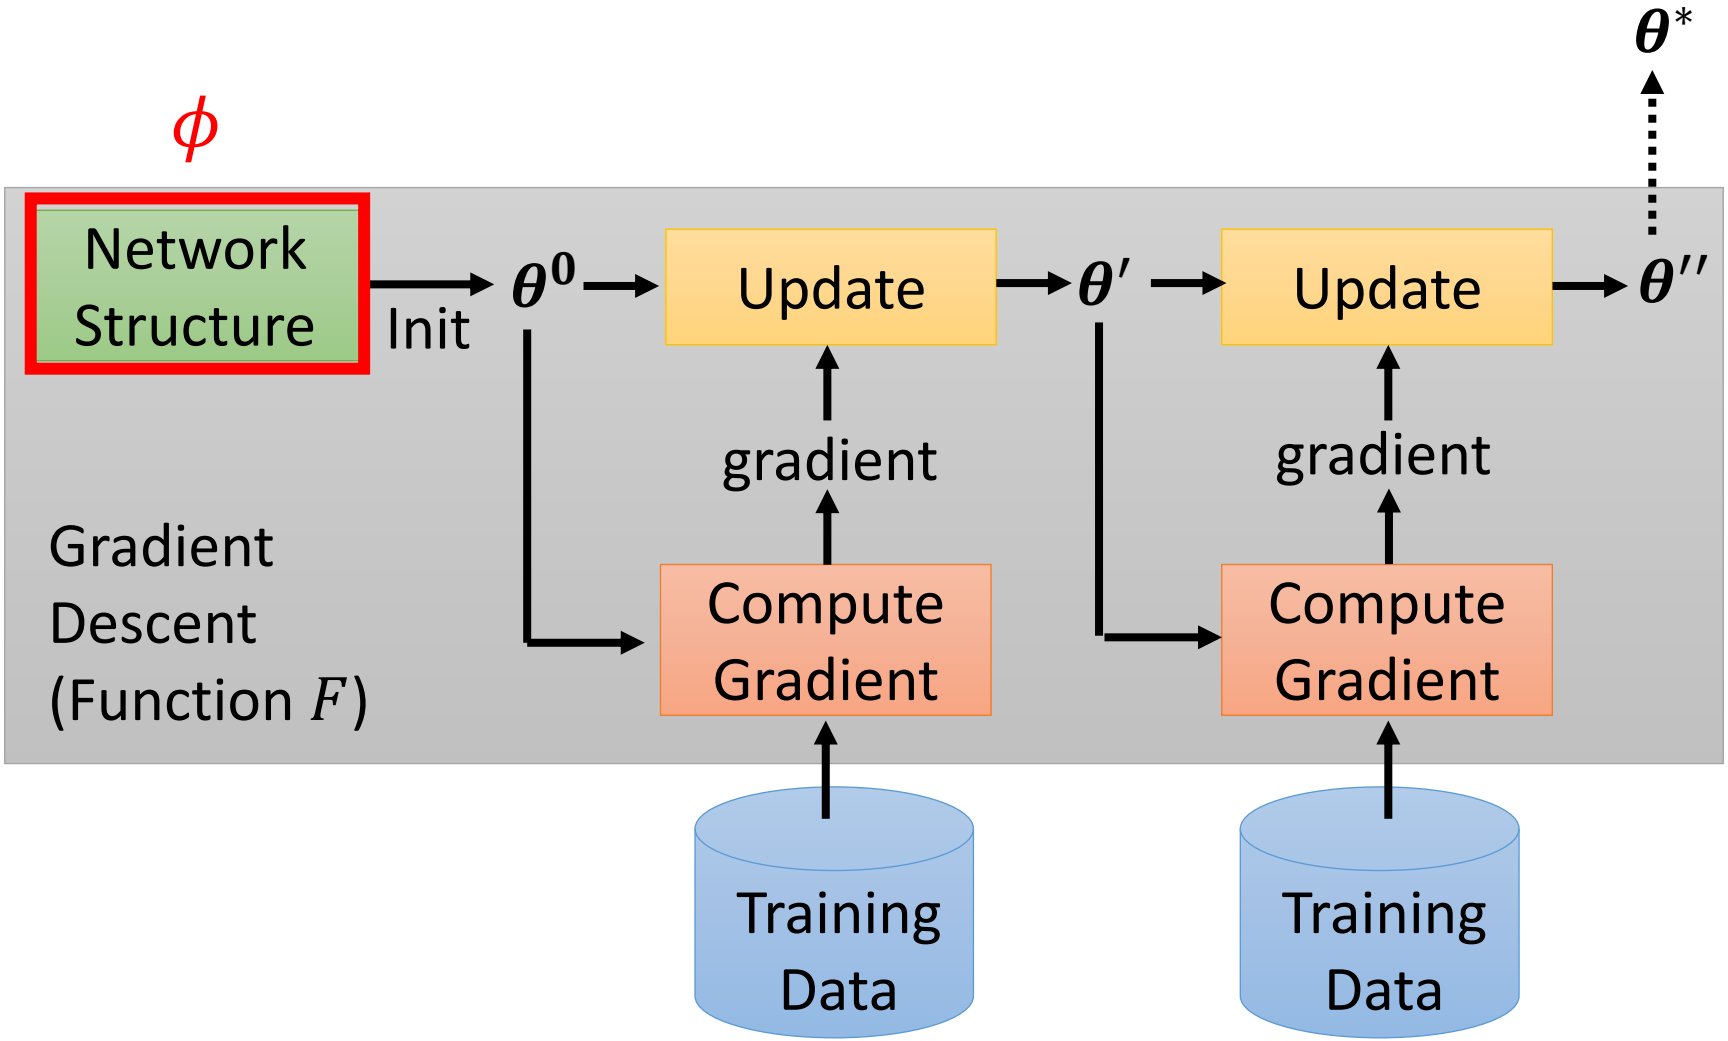

Network Architecture Search (NAS)

$$\hat{\phi} = arg\min_{\phi} L(\phi) \qquad \nabla_{\phi} L(\phi) = ?$$

ϕ: Network Architecture

- Reinforcement Learning
  - Barret Zoph, et al., Neural Architecture Search with Reinforcement Learning, ICLR 2017
  - Barret Zoph, et al., Learning Transferable Architectures for Scalable Image Recognition, CVPR, 2018
  - Hieu Pham, et al., Efficient Neural Architecture Search via Parameter Sharing, ICML, 2018

An agent uses a set of actions to determine the network architecture.

ϕ: the agent's parameters

$-L(\phi)$

Reward to be maximized

Network Architecture Search (NAS)

Across-task Training

Update ϕ to maximize reward $-L(\phi)$

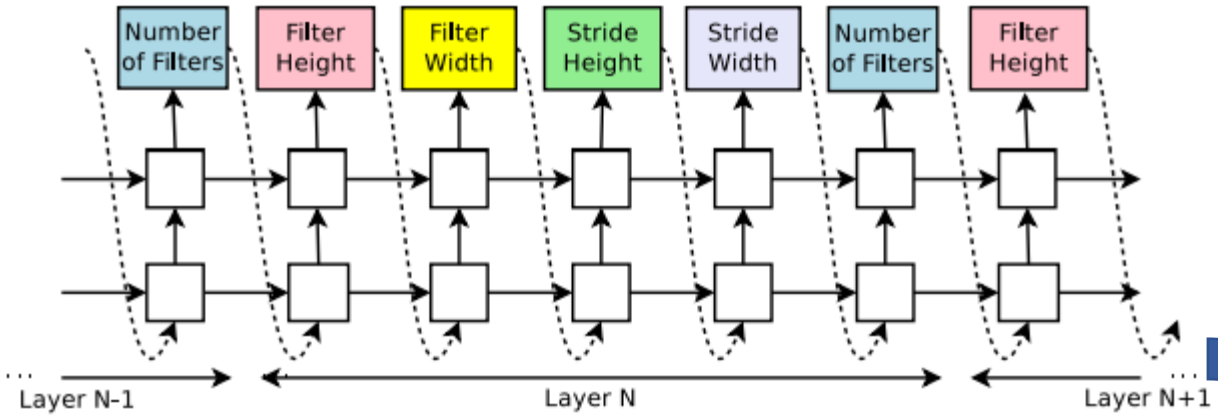

agent ϕ (RNN)

$-L(\phi)$

Accuracy of the network

form a network

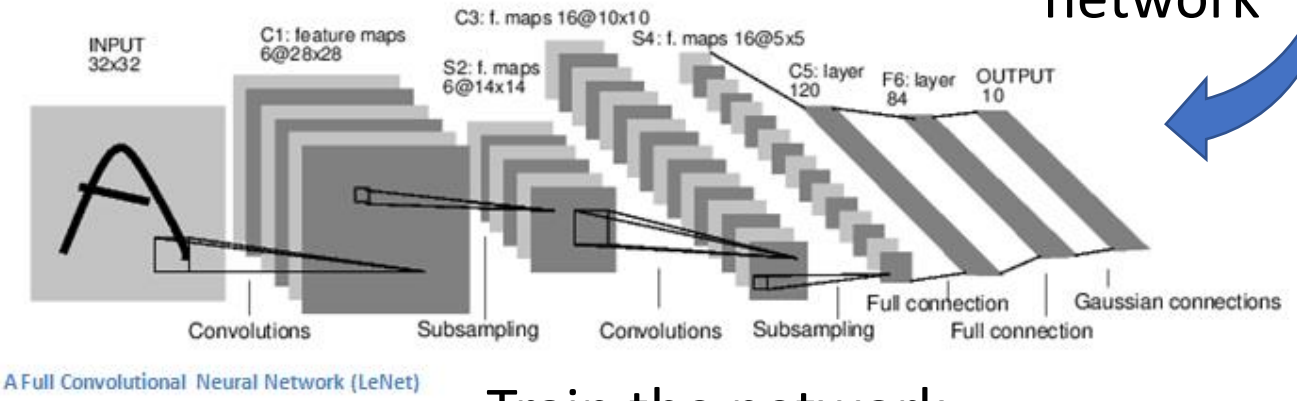

A Full Convolutional Neural Network (LeNet)

Train the network

Within-task Training

Network Architecture Search (NAS)

$$\hat{\phi} = arg\min_{\phi} L(\phi) \qquad \nabla_{\phi} L(\phi) = ?$$

ϕ: Network Architecture

- Reinforcement Learning
  - Barret Zoph, et al., Neural Architecture Search with Reinforcement Learning, ICLR 2017
  - Barret Zoph, et al., Learning Transferable Architectures for Scalable Image Recognition, CVPR, 2018
  - Hieu Pham, et al., Efficient Neural Architecture Search via Parameter Sharing, ICML, 2018
- Evolution Algorithm
  - Esteban Real, et al., Large-Scale Evolution of Image Classifiers, ICML 2017
  - Esteban Real, et al., Regularized Evolution for Image Classifier Architecture Search, AAAI, 2019
  - Hanxiao Liu, et al., Hierarchical Representations for Efficient Architecture Search, ICLR, 2018

Network Architecture Search (NAS)

$$\hat{\phi} = arg\min_{\phi} L(\phi) \qquad \nabla_{\phi} L(\phi) = ?$$

ϕ: Network Architecture

- DARTS — Hanxiao Liu, et al., DARTS: Differentiable Architecture Search, ICLR, 2019

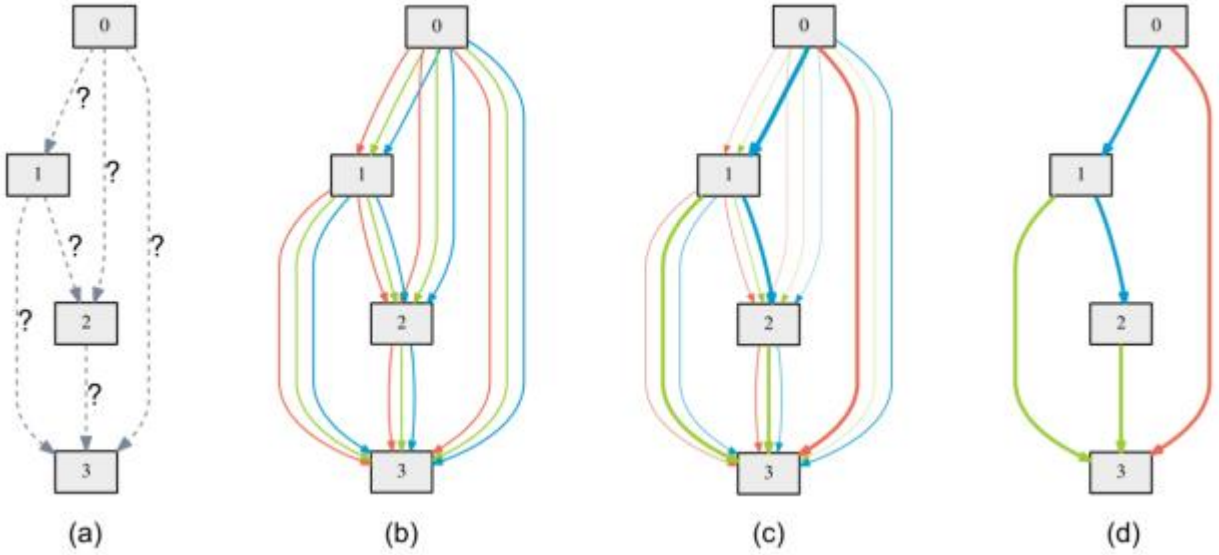

Data Processing?

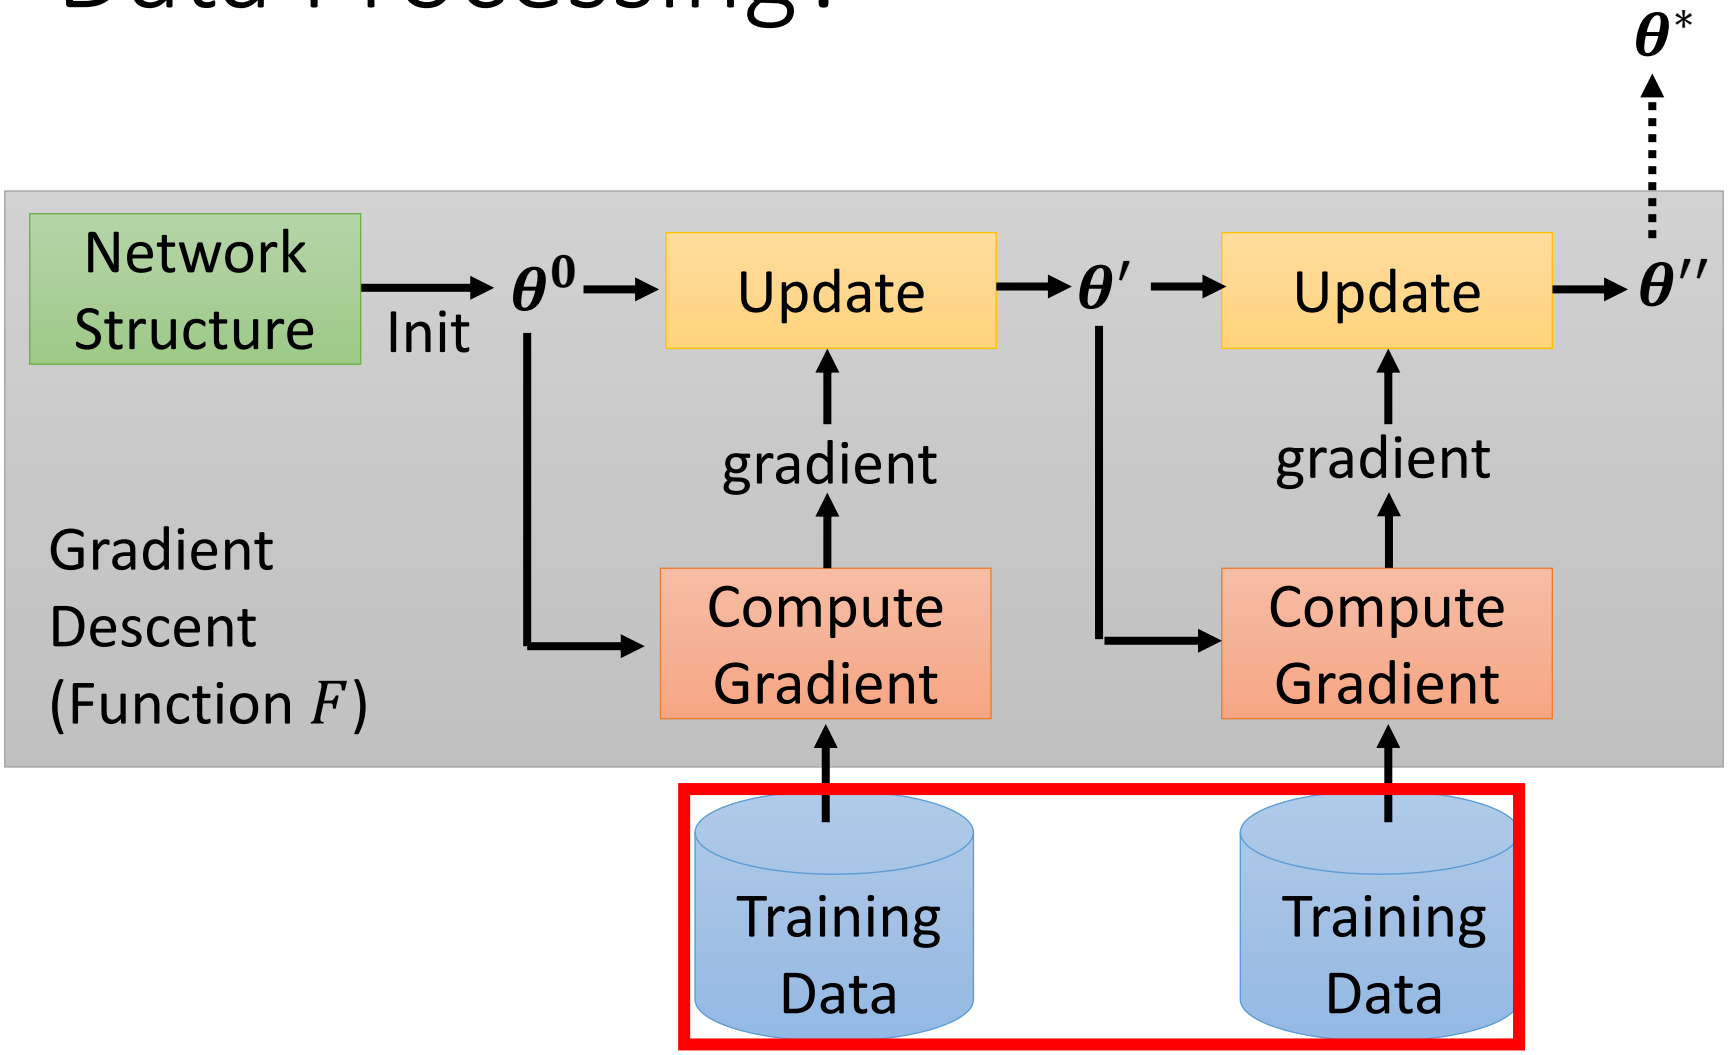

Data Augmentation

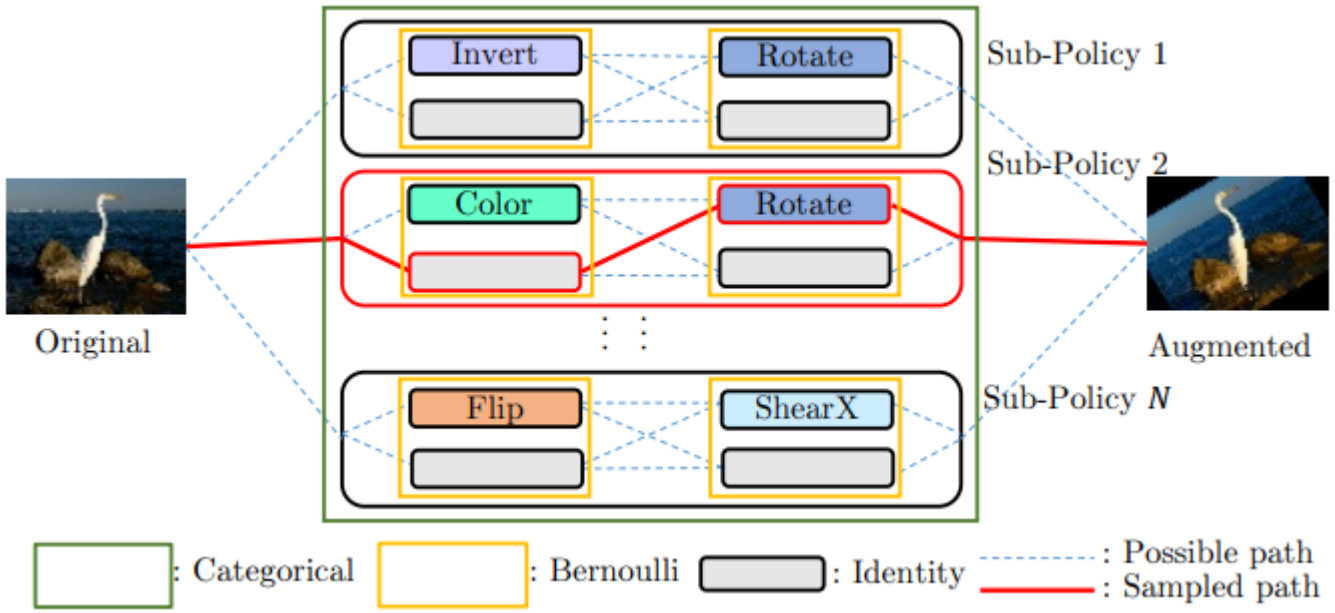

Yonggang Li, Guosheng Hu, Yongtao Wang, Timothy Hospedales, Neil M. Robertson, Yongxin Yang, DADA: Differentiable Automatic Data Augmentation, ECCV, 2020

Daniel Ho, Eric Liang, Ion Stoica, Pieter Abbeel, Xi Chen, Population Based Augmentation: Efficient Learning of Augmentation Policy Schedules, ICML, 2019

Ekin D. Cubuk, Barret Zoph, Dandelion Mane, Vijay Vasudevan, Quoc V. Le, AutoAugment: Learning Augmentation Policies from Data, CVPR, 2019

Sample Reweighting

- Give different samples different weights

Larger weights (focus on tough examples)?

Smaller weights (the labels are noisy)?

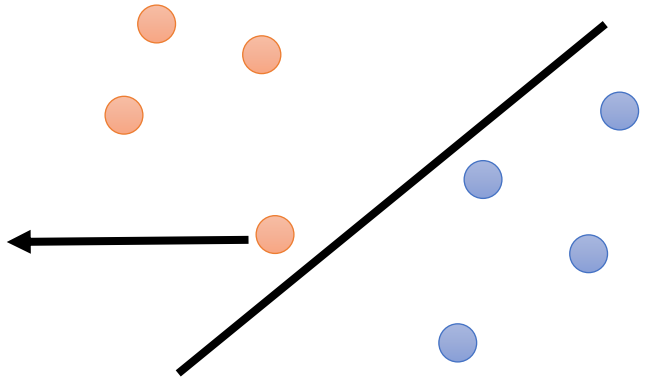

Sample Weighting Strategies → Learnable ϕ

Jun Shu, Qi Xie, Lixuan Yi, Qian Zhao, Sanping Zhou, Zongben Xu, Deyu Meng, Meta-Weight-Net: Learning an Explicit Mapping For Sample Weighting, NeurIPS, 2019

Mengye Ren, Wenyuan Zeng, Bin Yang, Raquel Urtasun, Learning to Reweight Examples for Robust Deep Learning, ICML, 2018

Beyond Gradient Descent

Andrei A. Rusu, Dushyant Rao, Jakub Sygnowski, Oriol Vinyals, Razvan Pascanu, Simon Osindero, Raia Hadsell,
Meta-Learning with Latent Embedding Optimization, ICLR, 2019

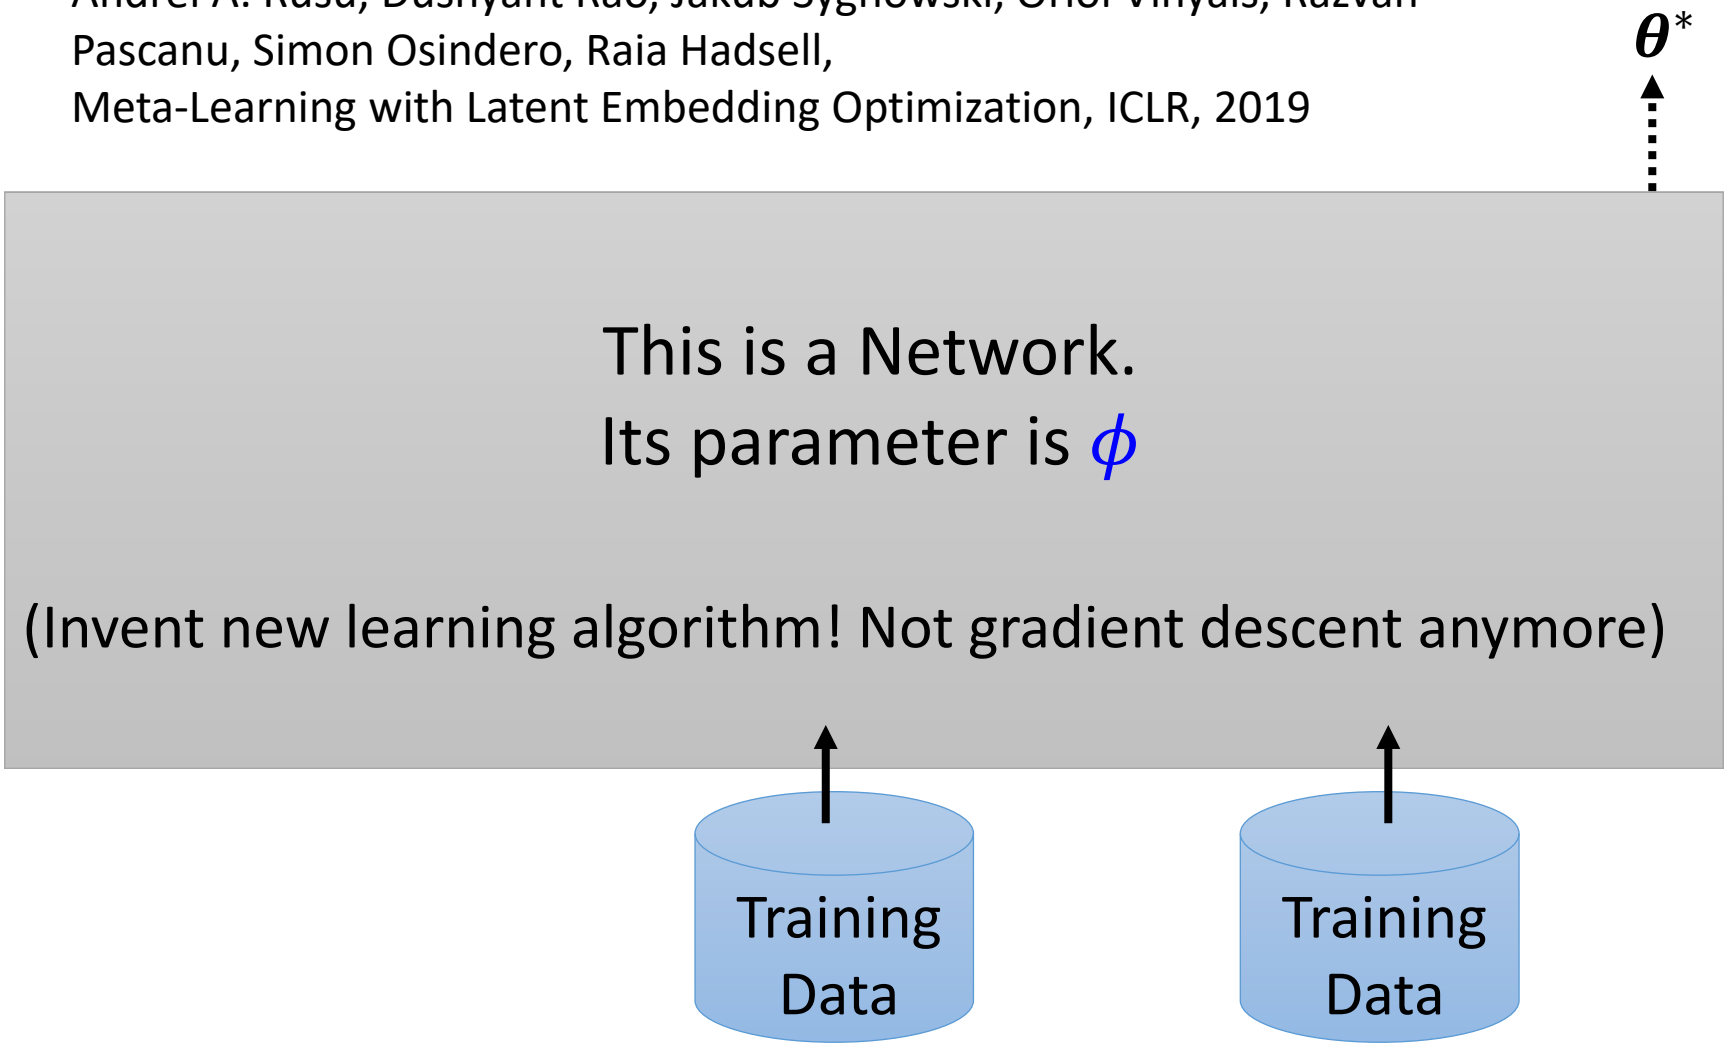

Until now ……

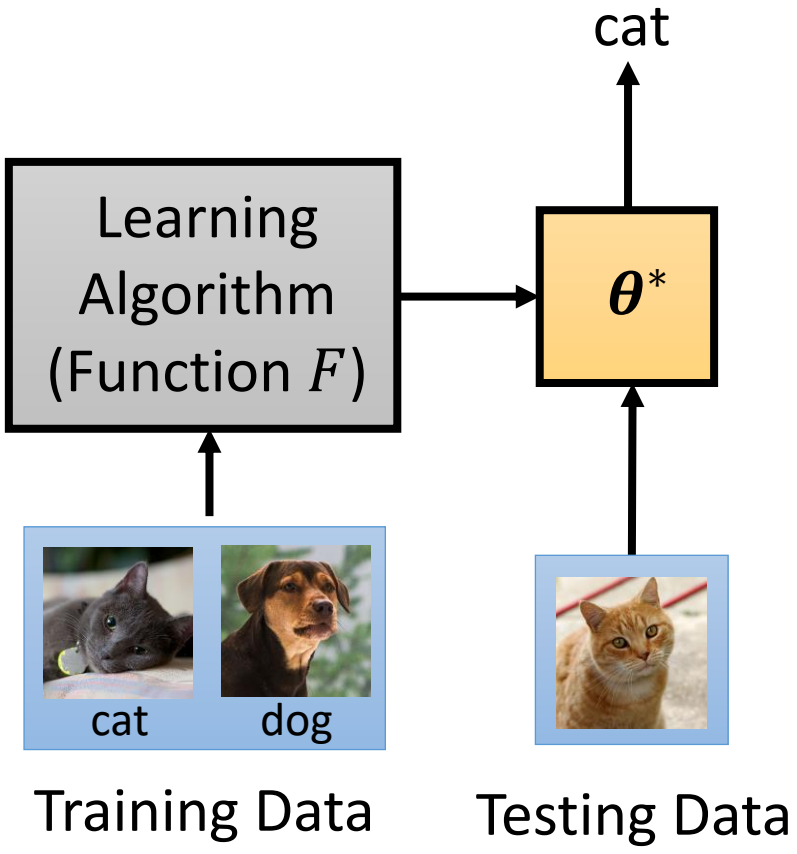

How about?

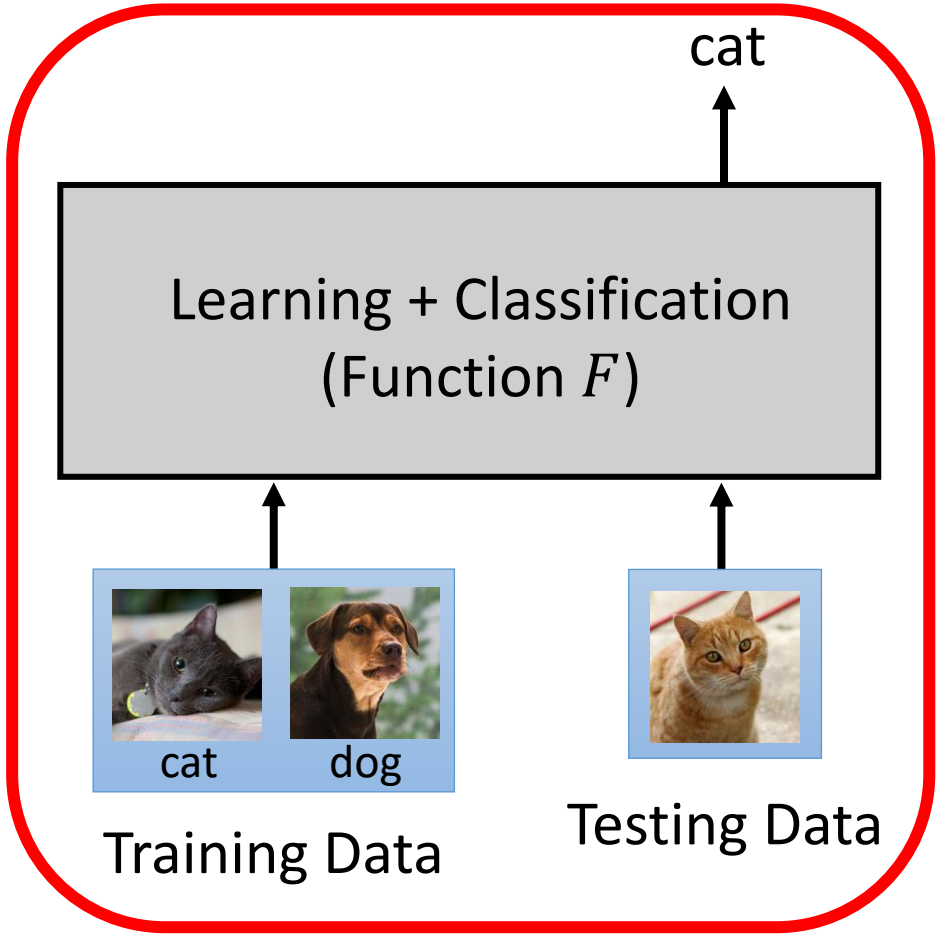

Learning to compare

(metric-based approach)

https://youtu.be/yyKaACh_j3M

https://youtu.be/scK2EIT7klw

https://youtu.be/semSxPP2Yzg

https://youtu.be/ePimv_k-H24

Applications

Few-shot Image Classification

- Each class only has a few images.

Class 1 Class 1 Class 2 Class 2 Class 3 Class 3 Which class?

3-ways 2-shot

- N-ways K-shot classification: In each task, there are N classes, each has K examples.
- In meta learning, you need to prepare many N-ways K-shot tasks as training and testing tasks.

Omniglot

https://github.com/brendenlake/omniglot

- 1623 characters
- Each has 20 examples

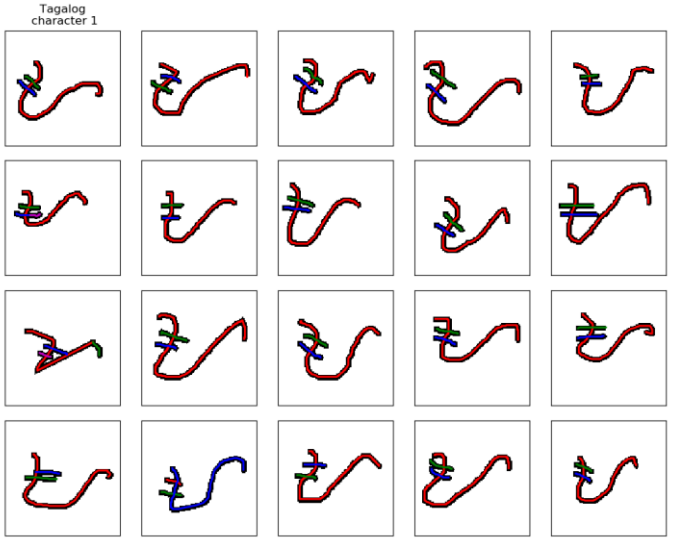

Omniglot

Demo:
https://openai.com/blog/reptile/

20 ways
1 shot

Each character represents a class

Testing set
(Query set)

Training set
(Support set)

- Split your characters into training and testing characters
  - Sample N training characters, sample K examples from each sampled characters → one training task
  - Sample N testing characters, sample K examples from each sampled characters → one testing task

	(A) **Learning to initialize**	(B) **Learning to compare**	(C) **Other**
Sound Event Detection	(Shi et al., 2020)	(Shimada et al., 2020a) (Chou et al., 2019) (Wang et al., 2020) (Shimada et al., 2020b) (Shi et al., 2020)	Network architecture search: (Li et al., 2020)
Keyword Spotting	(Chen et al., 2020a)	(Huh et al., 2020)	Net2Net: (Veniat et al., 2019) Network architecture search: (Mazzawi et al., 2019) Network architecture search: (Mo et al., 2020)
Text Classification	(Dou et al., 2019) (Bansal et al., 2019)	(Yu et al., 2018) (Tan et al., 2019) (Geng et al., 2019) (Sun et al., 2019)	Learning the learning algorithm: (Wu et al., 2019)
Voice Cloning			Learning the learning algorithm: (Chen et al., 2019b) (Serrà et al., 2019)
Sequence Labelng	(Wu et al., 2020)	(Hou et al., 2020)	
Machine Translation	(Gu et al., 2018) (Indurthi et al., 2020)		
Speech Recognition	(Hsu et al., 2020) (Klejch et al., 2019) (Winata et al., 2020a) (Winata et al., 2020b)		Learning to optimize: (Klejch et al., 2018) Network architecture search: (Chen et al., 2020b) (Baruwa et al., 2019)
Knowledge Graph	(Obamuyide and Vlachos, 2019) (Bose et al., 2019) (Lv et al., 2019) (Wang et al., 2019)	(Ye and Ling, 2019) (Chen et al., 2019a) (Xiong et al., 2018) (Gao et al., 2019)	
Dialogue / Chatbot	(Qian and Yu, 2019) (Madotto et al., 2019) (Mi et al., 2019)		Learning to optimize: (Chien and Lieow, 2019)
Parsing	(Guo et al., 2019) (Huang et al., 2018)		
Word Embedding	(Hu et al., 2019)	(Sun et al., 2018)	
Multi-model		(Eloff et al., 2019)	Learning the learning algorithm: (Surís et al., 2019)

http://speech.ee.ntu.edu.tw/~tlkagk/meta_learning_table.pdf

李宏毅 (Hung-yi Lee) · HYLEE | Machine Learning (2021)

HYLEE(2021)· 课程资料包 @ShowMeAI

视频
中英双语字幕

课件
一键打包下载

笔记
官方笔记翻译

代码
作业项目解析

视频 · B 站 [扫码或点击链接]
https://www.bilibili.com/video/BV1fM4y137M4

课件 & 代码 · 博客 [扫码或点击链接]
http://blog.showmeai.tech/ntu-hylee-ml

机器学习 Auto-encoder 生成式对抗网络 学习率
深度学习 卷积神经网络 GAN 自监督 自注意力机制
批次标准化 神经网络压缩 强化学习 元学习 Transformer

Awesome AI Courses Notes Cheatsheets 是 **ShowMeAI** 资料库的分支系列，覆盖最具知名度的 **TOP50+** 门 AI 课程，旨在为读者和学习者提供一整套高品质中文学习笔记和速查表。

点击课程名称，跳转至课程**资料包**页面，**一键下载**课程全部资料！

机器学习	深度学习	自然语言处理	计算机视觉
Stanford · CS229	Stanford · CS230	Stanford · CS224n	Stanford · CS231n
# Awesome AI Courses Notes Cheatsheets· 持续更新中			
知识图谱	**图机器学习**	**深度强化学习**	**自动驾驶**
Stanford · CS520	Stanford · CS224W	UCBerkeley · CS285	MIT · 6.S094

微信公众号

资料下载方式 2：扫码点击底部菜单栏
称为 **AI 内容创作者？** 回复 [添砖加瓦]